

# RSA<sup>®</sup>Conference2018

San Francisco | April 16 – 20 | Moscone Center

SESSION ID: EXP-W12

## THE RISE OF CONFIDENTIAL COMPUTING



#RSAC

**Mark Russinovich**

Azure CTO  
Microsoft



# Cloud Data Threats

## Customer cloud data concerns:



Malicious privileged  
admins or insiders



Hackers exploiting bugs  
in the Hypervisor/OS  
of cloud fabric



Third parties accessing  
it without customer  
consent

### CSA's Cloud Computing Top Threats


February 2016

#### Top Threat: Data Breaches



# Data Protection

## At rest



Encrypt inactive data when stored in blob storage, database, etc.

### Examples:

- Azure Storage Service Encryption for Data at Rest
- SQL Server Transparent Database Encryption (TDE)

## In transit




Encrypt data that is flowing between untrusted public or private networks

### Examples:

- HTTPS
- TLS

## In use



Protect/Encrypt data that is in use during computation

### Examples include:

- Trusted Execution Environments
- Homomorphic encryption

# Trusted Execution Environments (TEEs)

## Protected container:

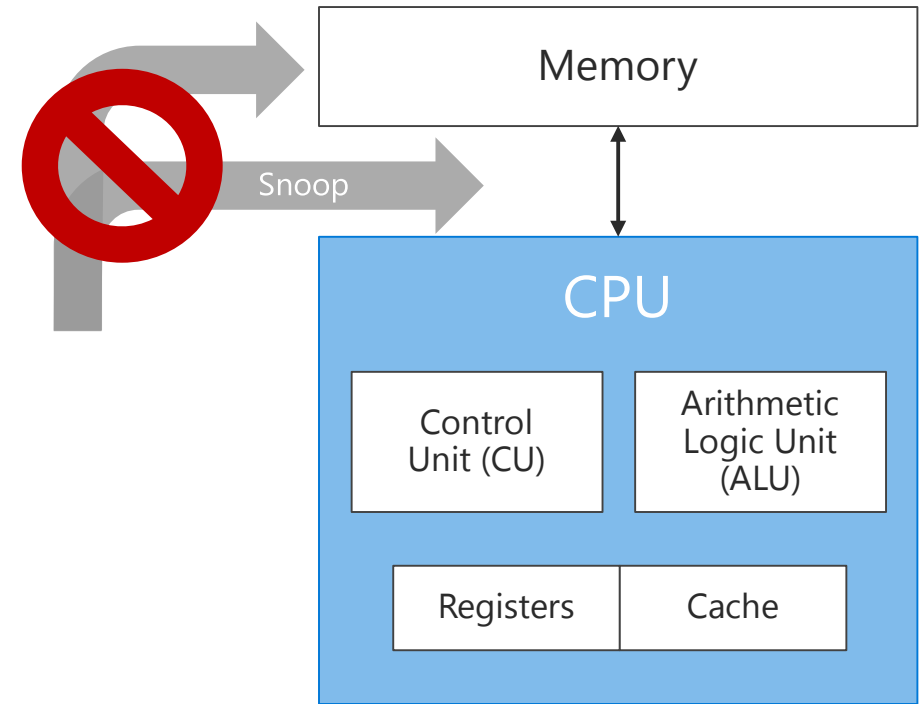
- Isolated portion of processor & memory
- Code & data cannot be viewed or modified from outside

Supports attestation: proving of identity both locally and remotely

Supports sealing: persisting secrets

## Examples:

- Intel SGX
- Virtualization Based Security (VBS) aka Virtual Secure Mode

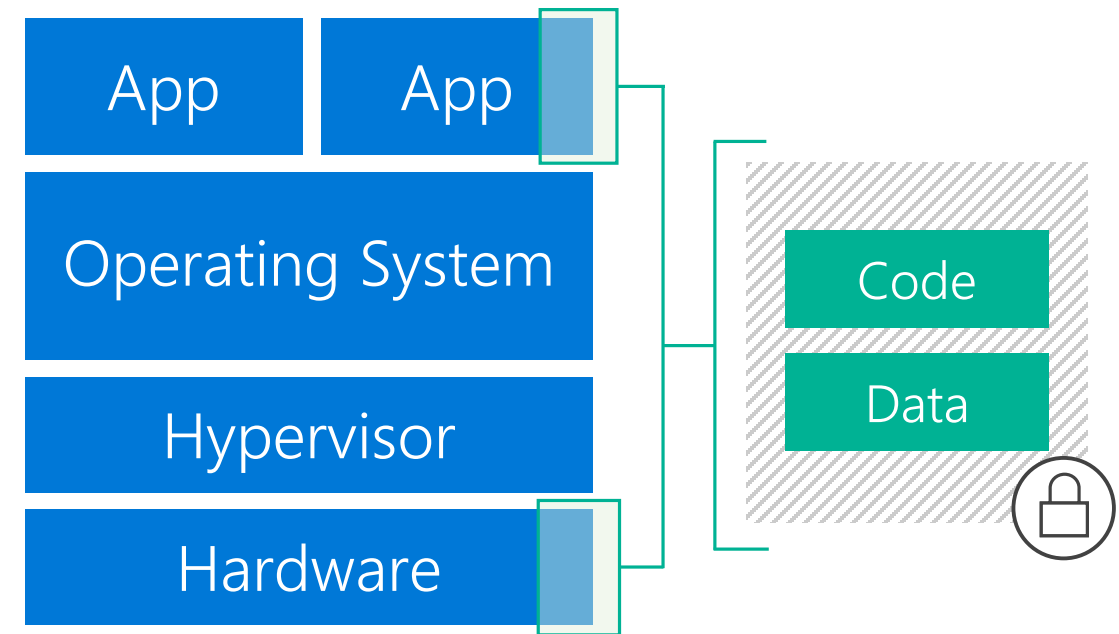


# Intel SGX (Software Guard Extensions)

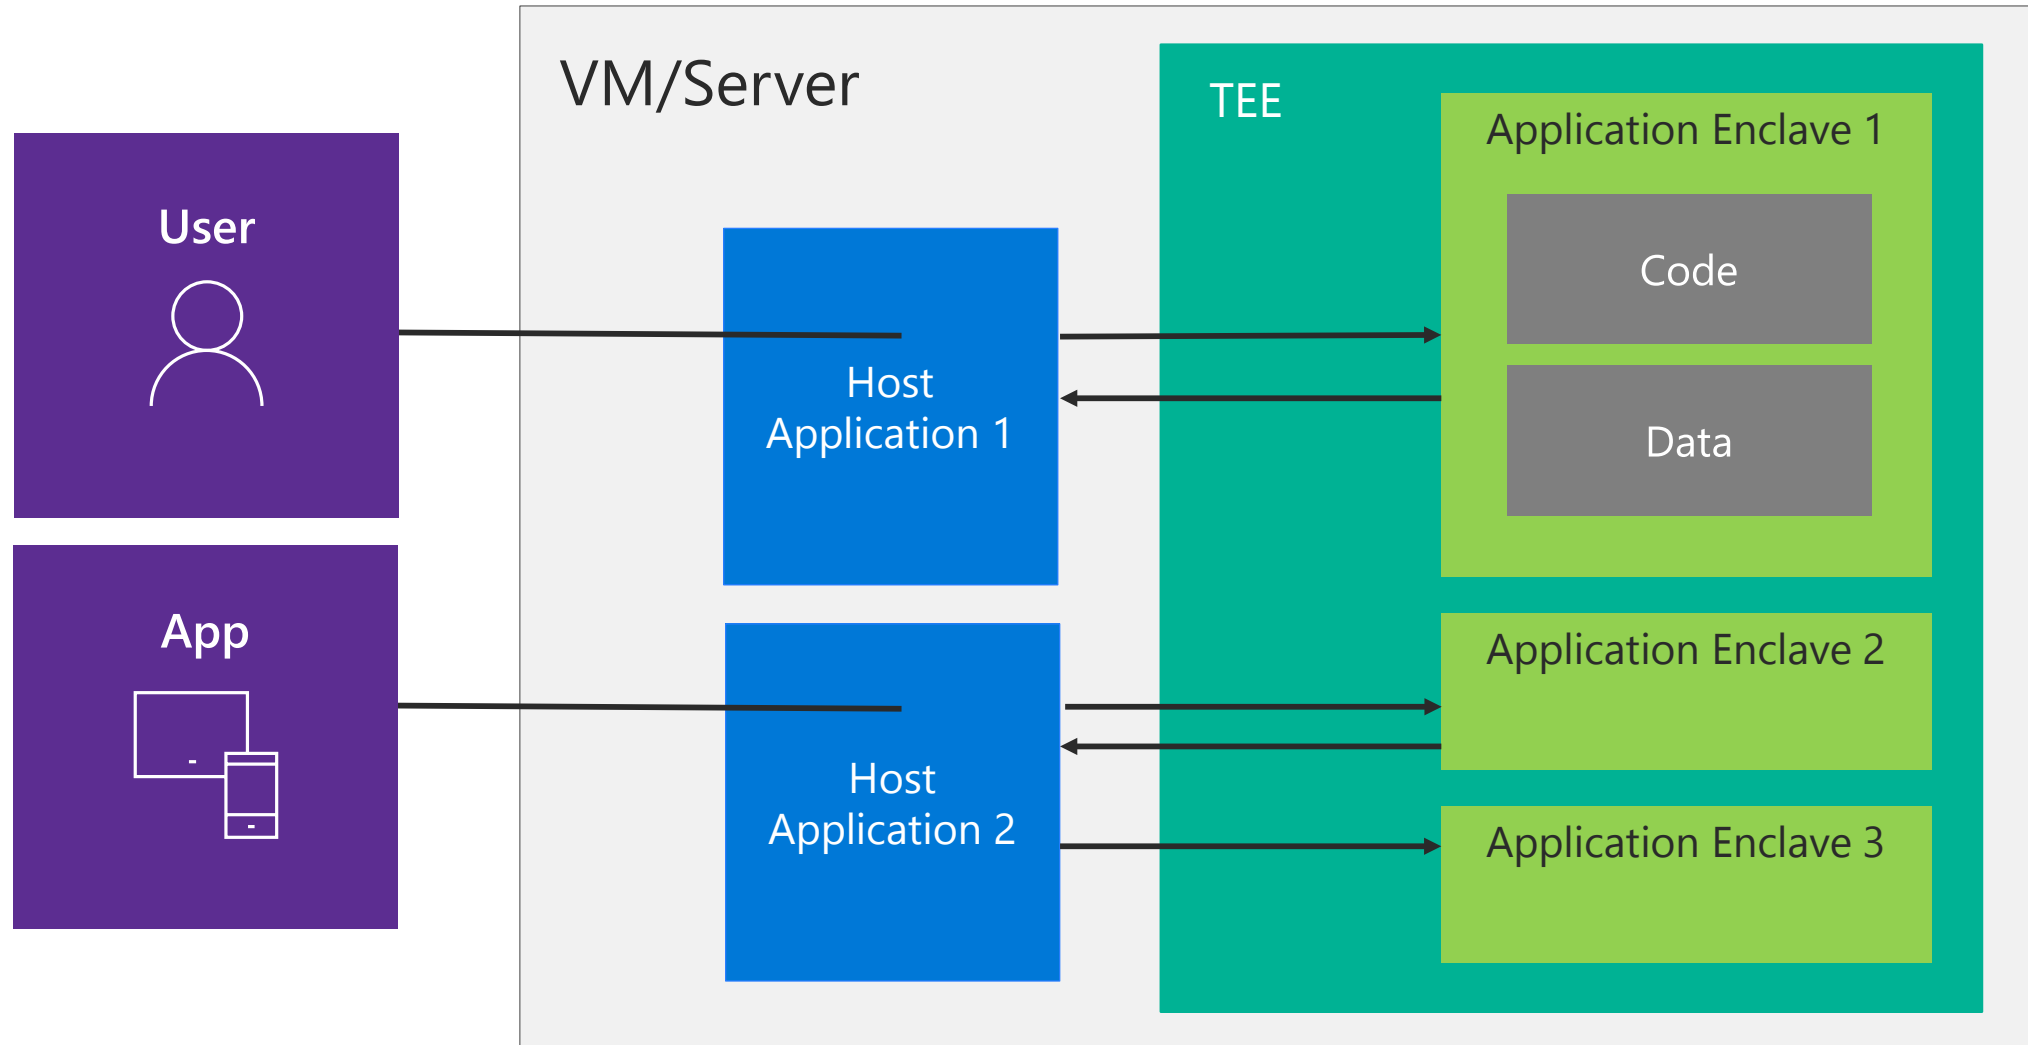
**SGX goal:** Minimize hardware attack surface

Instructions to set aside private regions ("protected containers") of code and data

Data is only in clear within protected memory and CPU



# TEE application architecture

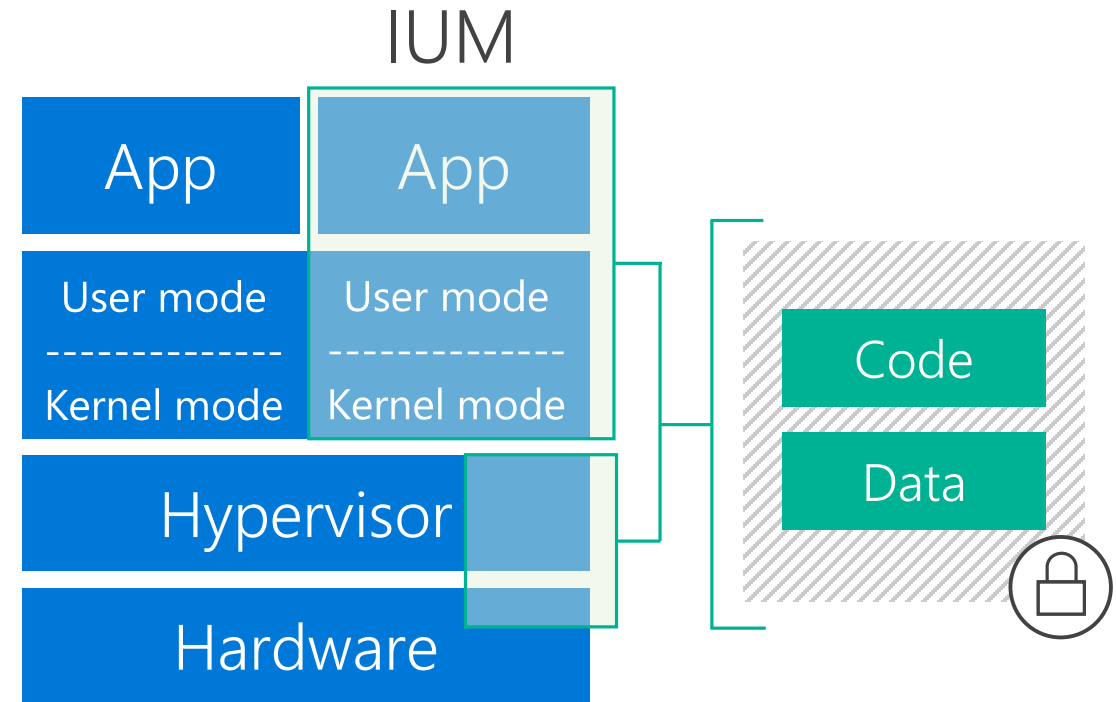


# Hyper-V Virtualization Based Security (VBS)\*

Virtual trust layers (VTLs) identify access levels

VTL 1 defines Isolated User Mode (IUM)

VBS prevents the OS, applications in VTL 0 and device drivers from accessing IUM (VTL 1)



\*AKA Virtual Secure Mode (VSM)

# TEEs compared to other secure hardware

	TPMs	HSMs	TEE (Intel SGX)
Hardware	Separate physical chip; embedded into motherboard	Separate external device	Built into CPU
Operations	Secure cryptographic operations	Secure cryptographic operations*	Secure “container” in which to run arbitrary code
Attestation	Yes	Yes	Yes
Sealing	Yes	Yes	Yes
Use case	Local system integrity <ul style="list-style-type: none"><li>• Validate properties at boot</li><li>• Host measurement</li><li>• Platform device authentication</li></ul>	Key management system <ul style="list-style-type: none"><li>• Generate &amp; manage keys</li><li>• Key exchange</li><li>• Encryption</li></ul>	Generalized compute
Example application	BitLocker Disk Encryption	KeyVault	SQL Always Encrypted

\*Some permit arbitrary code, but not optimized for general purpose compute



# Common TEE application patterns



Protect data confidentiality and integrity on remote machine

Protect sensitive algorithmic IP (e.g. financial trade algorithms)

Code access security, including remote clients

Create a trusted network of nodes among a set of untrusted parties

Centrally combine different data sources for a better algorithmic outcome without loss of confidentiality

Protect communication with other secure device endpoints (e.g. local peer to remote HSM)

Secure licensing and DRM

# Azure Confidential Computing

# Confidential cloud



**Data is fully in the control of the customer** regardless of whether in rest, transit, or use even though the infrastructure is not



The cloud platform provider is outside the **trusted compute base**



**Code** running in cloud is protected and verified **by the customer**

# Azure and confidential computing



Working with silicon partners to enable Confidential Computing

Building software to deploy, manage, and develop secure TEE applications on Azure

Designing and developing services to support attestation in the cloud


Enabling confidential PaaS and SaaS services

# The ACC development environment




## Universal

Generalize enclave application model to minimize hardware/software specific concepts



## Pluggable

Componentization to support desired runtimes and crypto libraries



## Standardized

Remove hardware vendor specific signing and verification requirements



## Multi-platform

Design with all software platforms, Windows and Linux, in mind



## Compatible

Easier enablement of redistributable applications

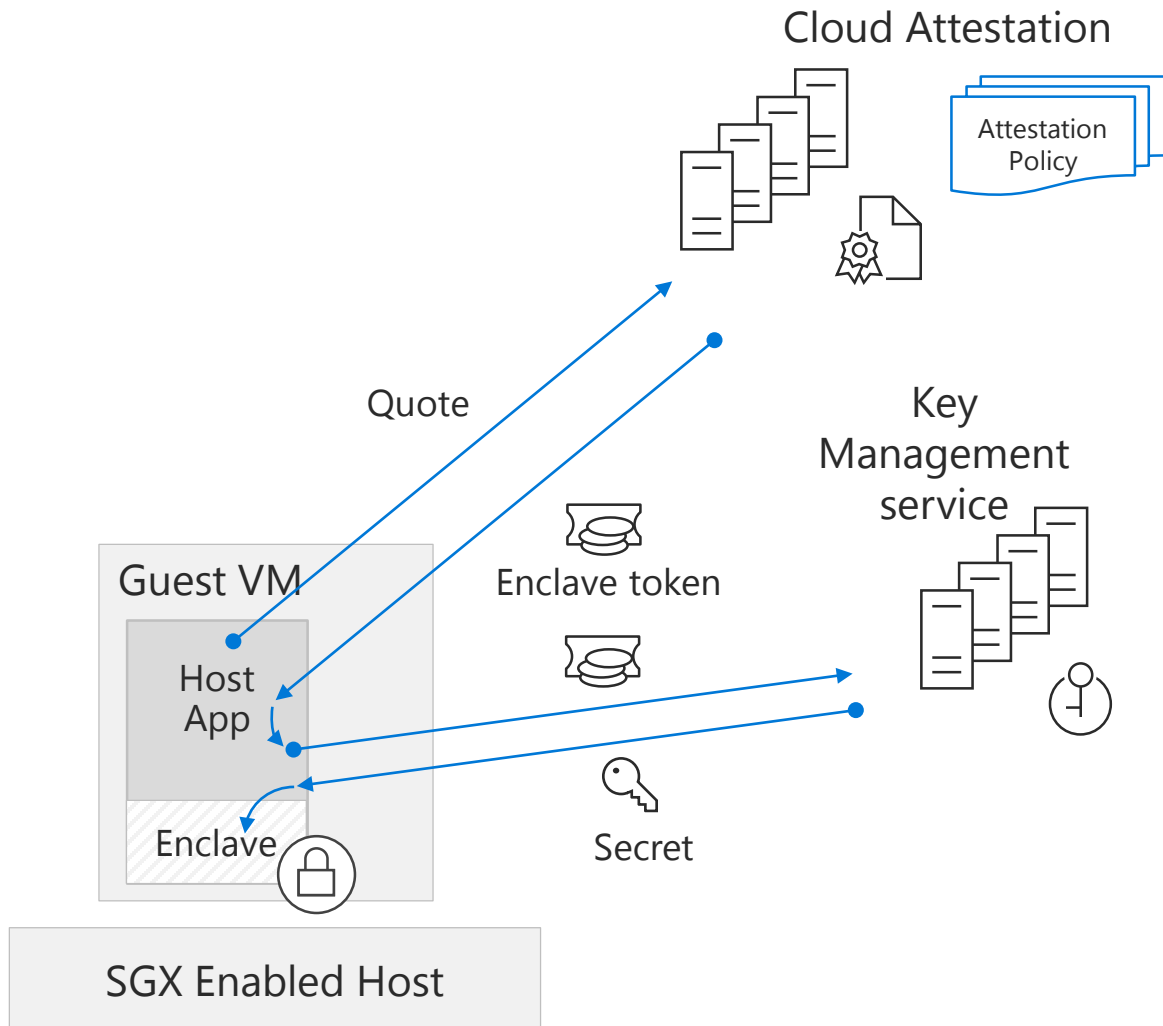


## Open

Open source and a standard for secure enclave-based application



# Universal cloud attestation



1. Quote provide proofs:
  - Code runs in genuine SGX enclave
  - Enclave version and owner is as expected
  - Arbitrary enclave supplied data is as expected
2. Attestation service attests to hardware, rooted in Intel chain of trust, and issues token
3. Cloud service (e.g. AKV) is presented with token with certs chained to CPU
4. Token is used to release secrets to the application enclave

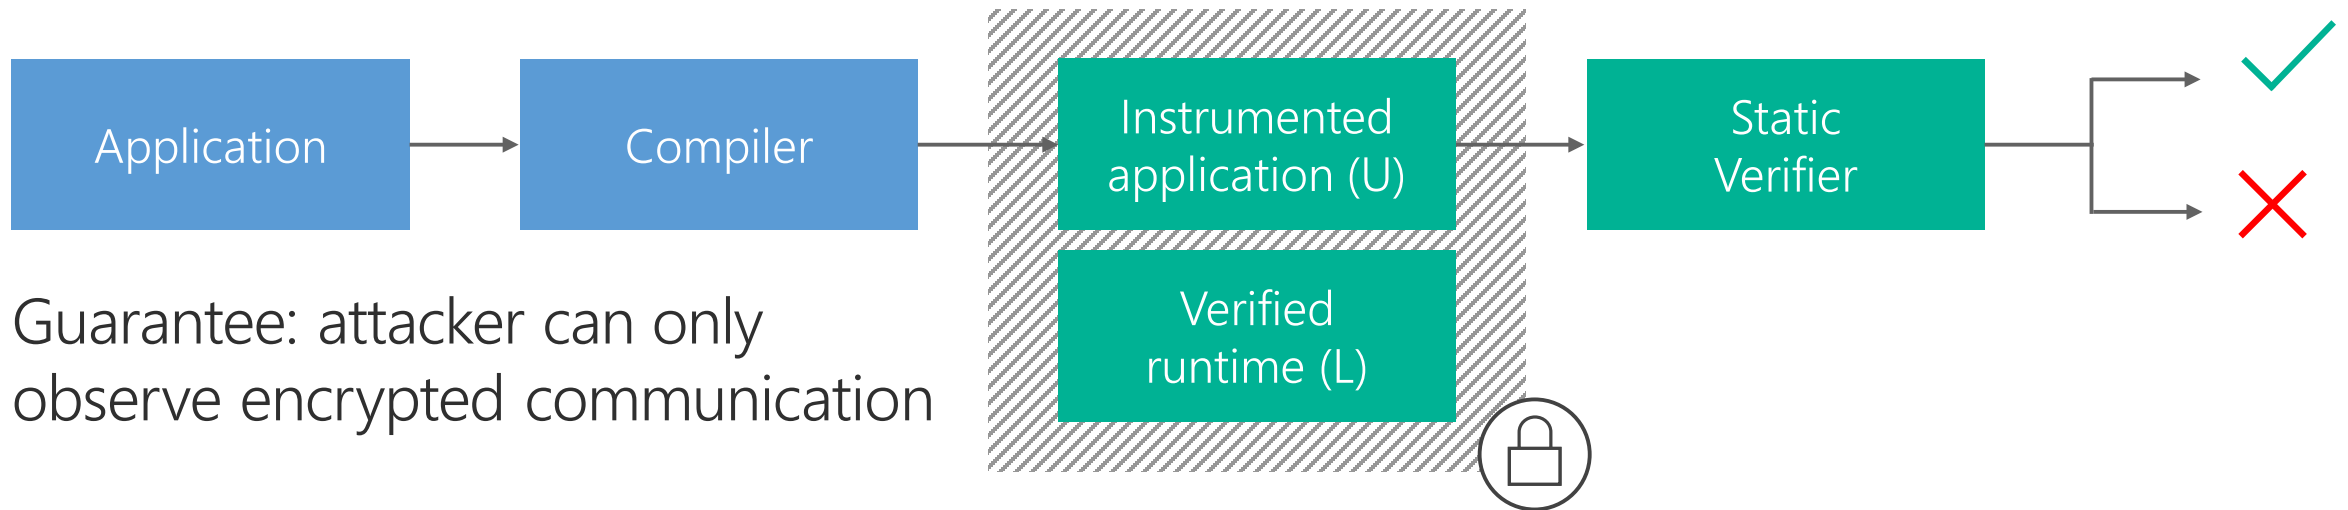
# Preventing direct information leaks



**Problem:** code in enclaves may unintentionally write secrets out



**Solution:** use a compiler that instruments memory accesses & verify that instrumented binary does not leak secrets



Guarantee: attacker can only observe encrypted communication

# Preventing indirect information leaks

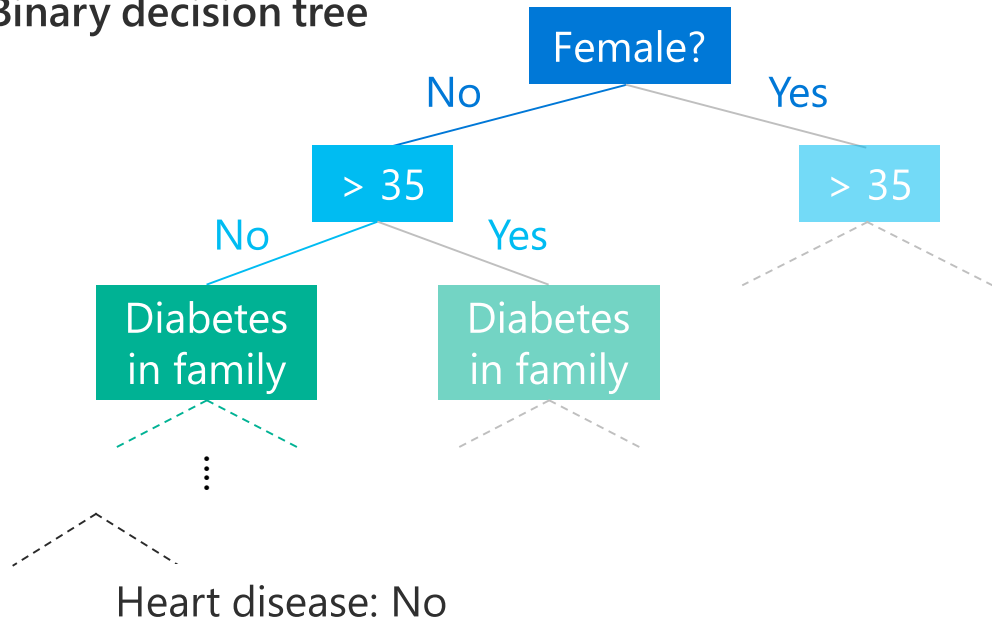


**Problem:** memory/disk access patterns may leak information




**Solution:** use compiler and hardened libraries that prevent leaks with data oblivious primitives

Binary decision tree



Memory



# Demo:

# Oblivious computing

# Example confidential computing scenarios






Always encrypted storage  
with  
SQL Server



Enabling scalable  
and confidential  
blockchain networks  
with Coco Framework



Financial data  
processing

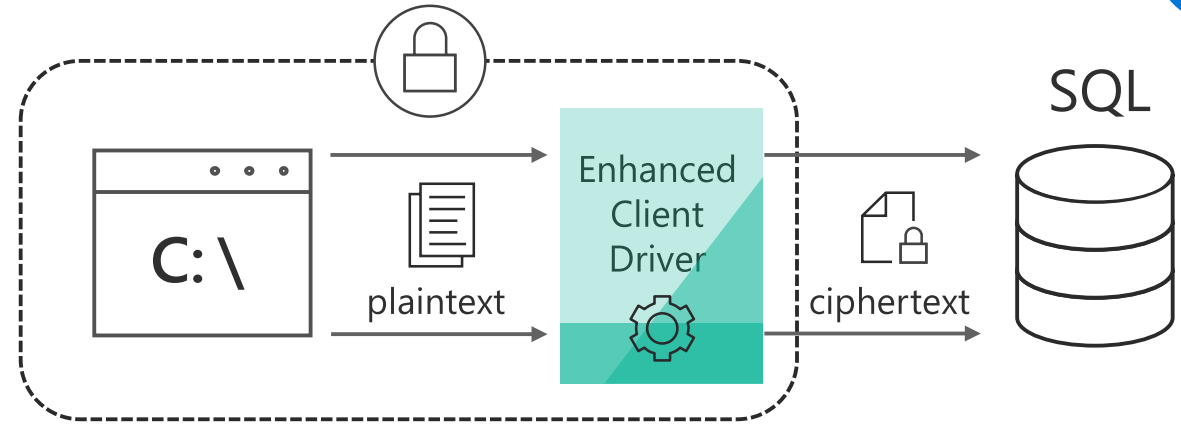


Secure multi-party  
machine learning

# SQL Always Encrypted

Protects sensitive data **in use** from high-privileged yet unauthorized SQL users both on-premises and in the cloud

Current GA version in SQL Server 2016/17 and Azure SQL DB



## Client side Encryption

Client-side encryption of sensitive data using keys that are **never** given to the database system

## Encryption Transparency

Client driver transparently encrypts query parameters and decrypts encrypted results

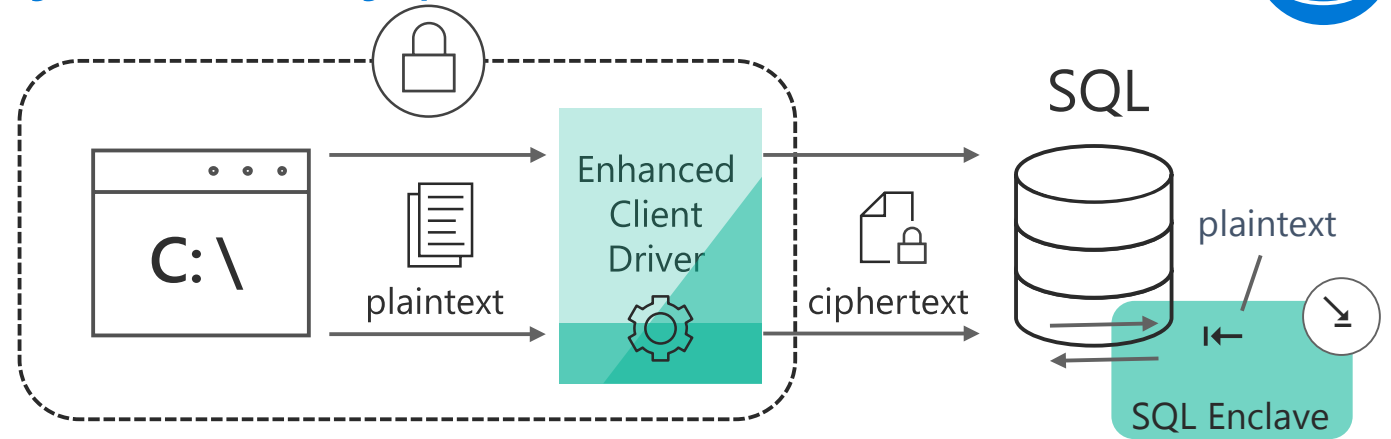
## Queries on Encrypted Data

Support for equality comparison, including join, group by and distinct operators via deterministic encryption

# Confidential SQL Always Encrypted



Protects sensitive data in use **while preserving rich queries and providing in-place encryption**



## Secure computations inside SQL Enclave

SQL Server Engine delegates operations on encrypted to the SQL Enclave, where the data can be safely decrypted and processed

## Rich Queries

pattern matching (LIKE), range queries (<, >, etc.), sorting, type conversions, support for non-bin2 collation, and more

## In-place Encryption

SQL Enclave can perform initial data encryption and key rotation, without moving the data out of the database


# Coco Framework: Confidential Consortium Blockchain Framework



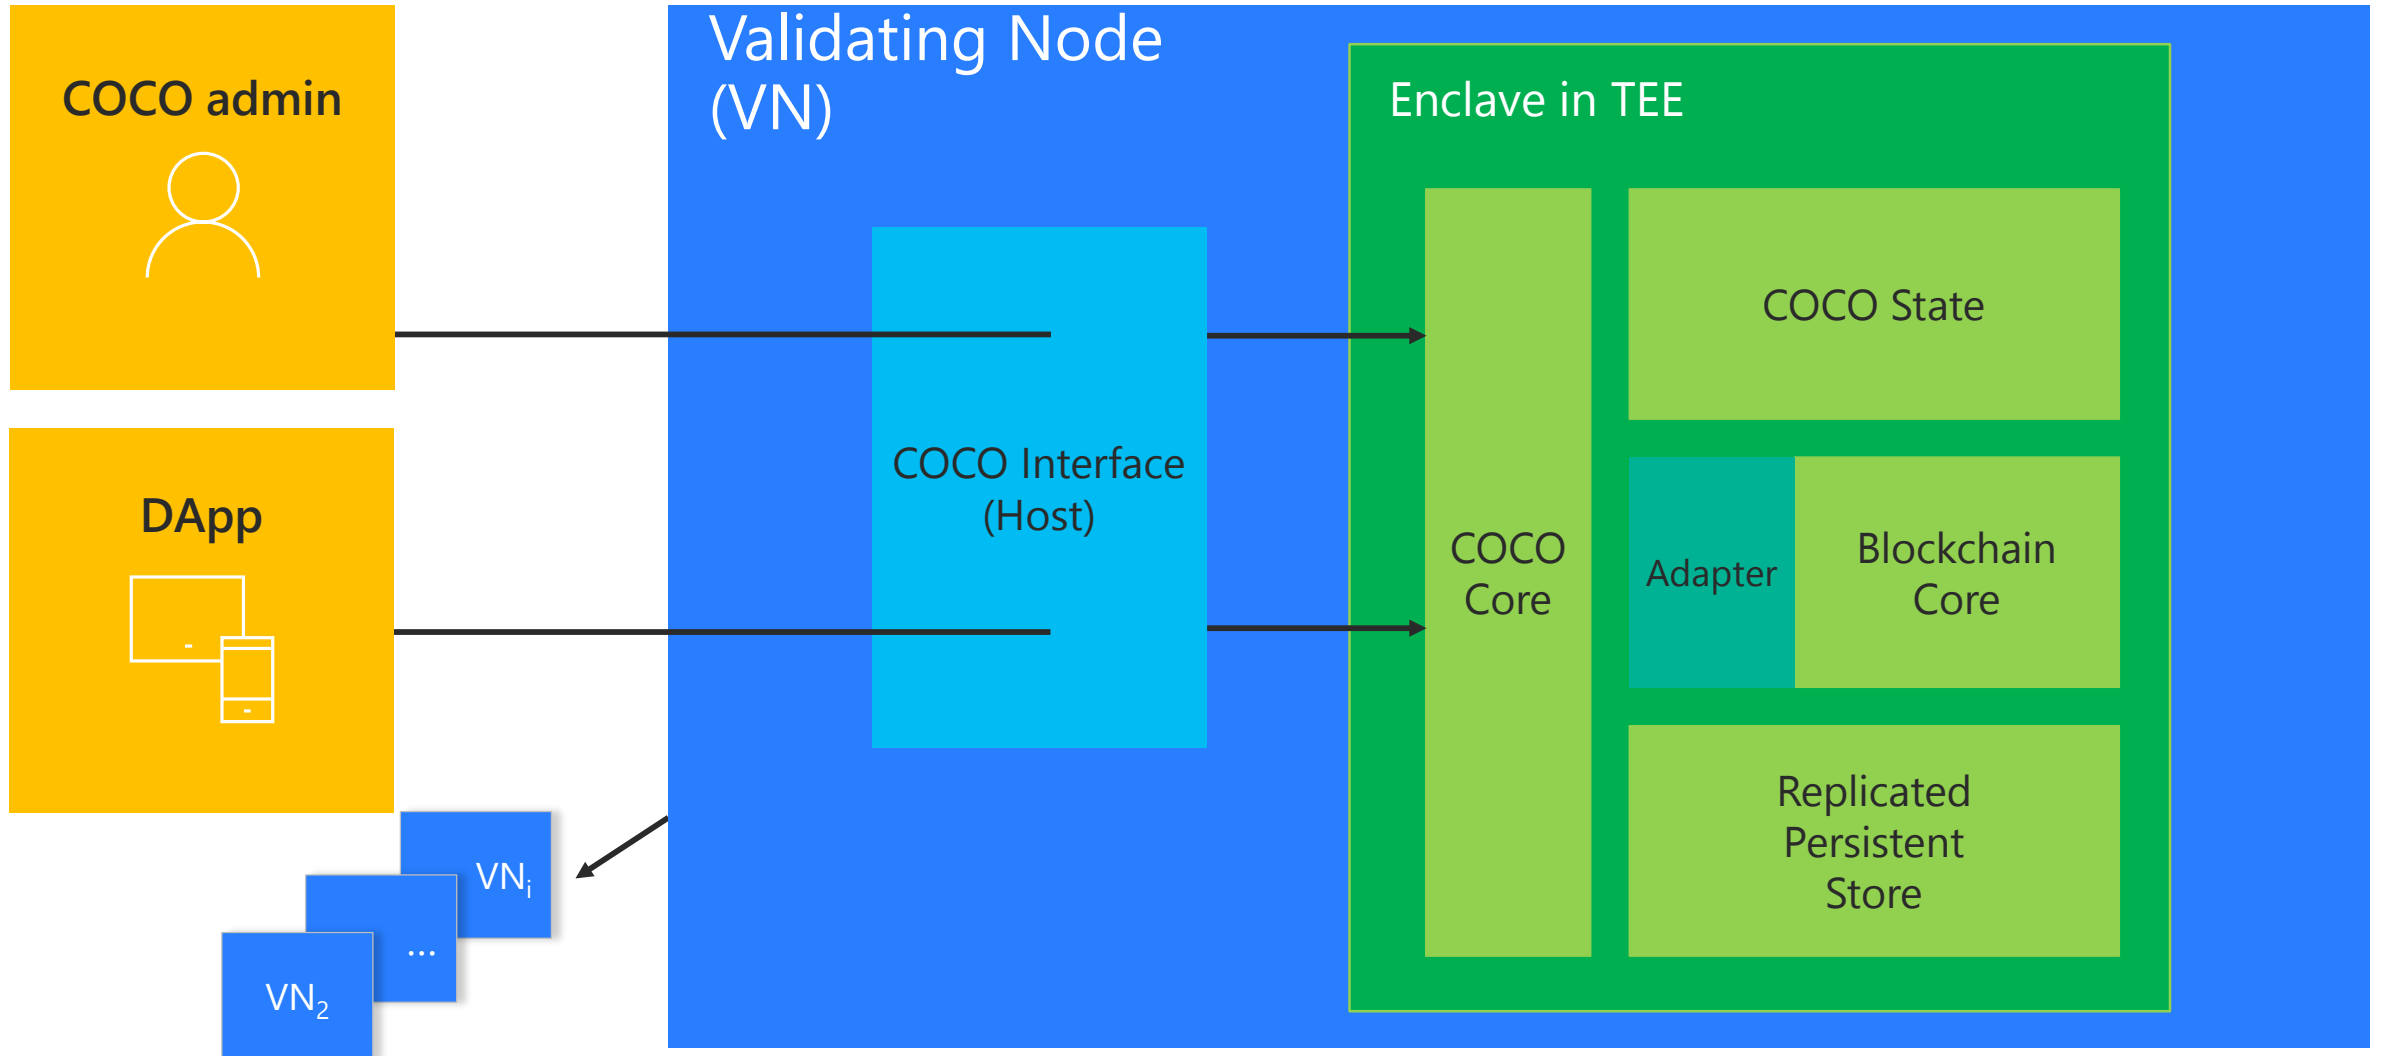
Open-source framework that enables high-throughput (~100x), fine-grained confidentiality, and consortium governance for blockchain

Creates a trusted network of physical nodes on which to run a distributed ledger, providing secure, reliable components for the protocol to use

Through the use of TEEs able to simplify consensus and transaction processing



# Coco Framework architecture






# Coco Framework: Confidentiality model for Ethereum



Coco FX disables access to transaction and block level information

Ethereum smart contract enforces access control rules by verifying address of caller



# Smart contract access control



```
98 // Only admins can authorize other users or admins
99
100 function addAuthorizedReader(address a) onlyAdmins
101 {
102     allowedReaders[a] = true;
103     userAuthorized(a);
104 }
105
106 function addAuthorizedWriter(address a) onlyAdmins
107 {
108     allowedWriters[a] = true;
109     userAuthorized(a);
110 }
111
112 function addAdmin(address a) onlyAdmins
113 {
114     admins[a] = true;
115     userAuthorized(a);
116 }
117
118 function removeAuthorizedReader(address a) onlyAdmins
119 {
120     allowedReaders[a] = false;
121     userForbidden(a);
122 }
123
124 function removeAuthorizedWriter(address a) onlyAdmins
125 {
126     allowedWriters[a] = false;
127     userForbidden(a);
128 }
129
130 function removeAdmin(address a) onlyAdmins
131 {
132     admins[a] = false;
133     userForbidden(a);
134 }
135
136 }
```

# Smart contract access control (cont)



```
1 pragma solidity ^0.4.0;
2
3 /* A sample contract that restricts access to a map of values. A user is represented by
4  * an Ethereum address they own.
5  */
6
7 contract Vault
8 {
9     // Each access is logged and visible only to authorized users
10    event vaultReadAccess(address byWhom);
11    event vaultWriteAccess(address byWhom);
12    event userAuthorized(address who);
13    event userForbidden(address who);
14
15    /* The map that is access restricted and the lists of users that are permitted to
16     * read and write. Note that these members are not public and can only be accessed
17     * by code in this contract
18     */
19    mapping (bytes32 => string) secretsMap;
20
21    mapping (address => bool) allowedReaders;
22    mapping (address => bool) allowedWriters;
23    mapping (address => bool) allowedEventViewers;
24
25    // Admins can authorize users and add other admins
26    mapping (address => bool) admins;
27
28    // Check
29    function check(bytes32 hash, uint8 v, bytes32 r, bytes32 s, bytes32 topic) returns (bool){
30        address sender = ecrecover(hash, v, r, s);
31
32        bytes32 event1Topic = keccak256("vaultReadAccess(address)");
33        bytes32 event2Topic = keccak256("vaultWriteAccess(address)");
34        bytes32 event3Topic = keccak256("userAuthorized(address)");
35        bytes32 event4Topic = keccak256("userForbidden(address)");
36
37        if (topic == event1Topic){
38            //`allowedReaders` can access event: `vaultReadAccess`
39            return allowedReaders[sender];
40        }else if (topic == event2Topic){
41            //`allowedWriters` can access event: `vaultWriteAccess`
42            return allowedWriters[sender];
43        }else if (topic == event3Topic || topic == event4Topic){
44            //`admins` can access event: `userAuthorized`, `userForbidden`
45            return admins[sender];
46        }else {
47            throw ;
48        }
49    }
50
51    // Contract is created with an initial set of authorized users who can read, write and view events
52    function Vault(address[] initialAdmins)
53    {
54        for (uint i = 0; i< initialAdmins.length;i++){
55            admins[initialAdmins[i]] = true;
56        }
57    }
58
59    // A read call made via eth_call is not signed so we require the caller supply an
60    // ECDSA signature to authenticate themselves
61    modifier onlyAllowedReaders(bytes32 hash, uint8 v, bytes32 r, bytes32 s)
62    {
63        address reader = ecrecover(hash, v, r, s);
64        if (!allowedReaders[reader])
65            throw;
66        _;
67    }
68
69    modifier onlyAllowedWriters
70    {
71        if (!allowedWriters[msg.sender])
72            throw;
73        _;
74    }
75
76    modifier onlyAdmins
77    {
78        if (!admins[msg.sender])
79            throw;
80        _;
81    }
82
83    // Modifier onlyAllowedReaders ensures only users authorized to read can call this function
84    function read(bytes32 hash, uint8 v, bytes32 r, bytes32 s, bytes24 key) onlyAllowedReaders(hash, v, r, s) constant returns(string)
85    {
86        vaultReadAccess(msg.sender); // Generate an event for notification
87        return secretsMap[key];
88    }
89
90    // Modifier onlyAllowedWriters ensures only users authorized to write can call this function
91    function write(bytes32 key, string value) onlyAllowedWriters
92    {
93        secretsMap[key] = value;
94        vaultWriteAccess(msg.sender); // Generate an event for notification
95    }
96
97 }
```

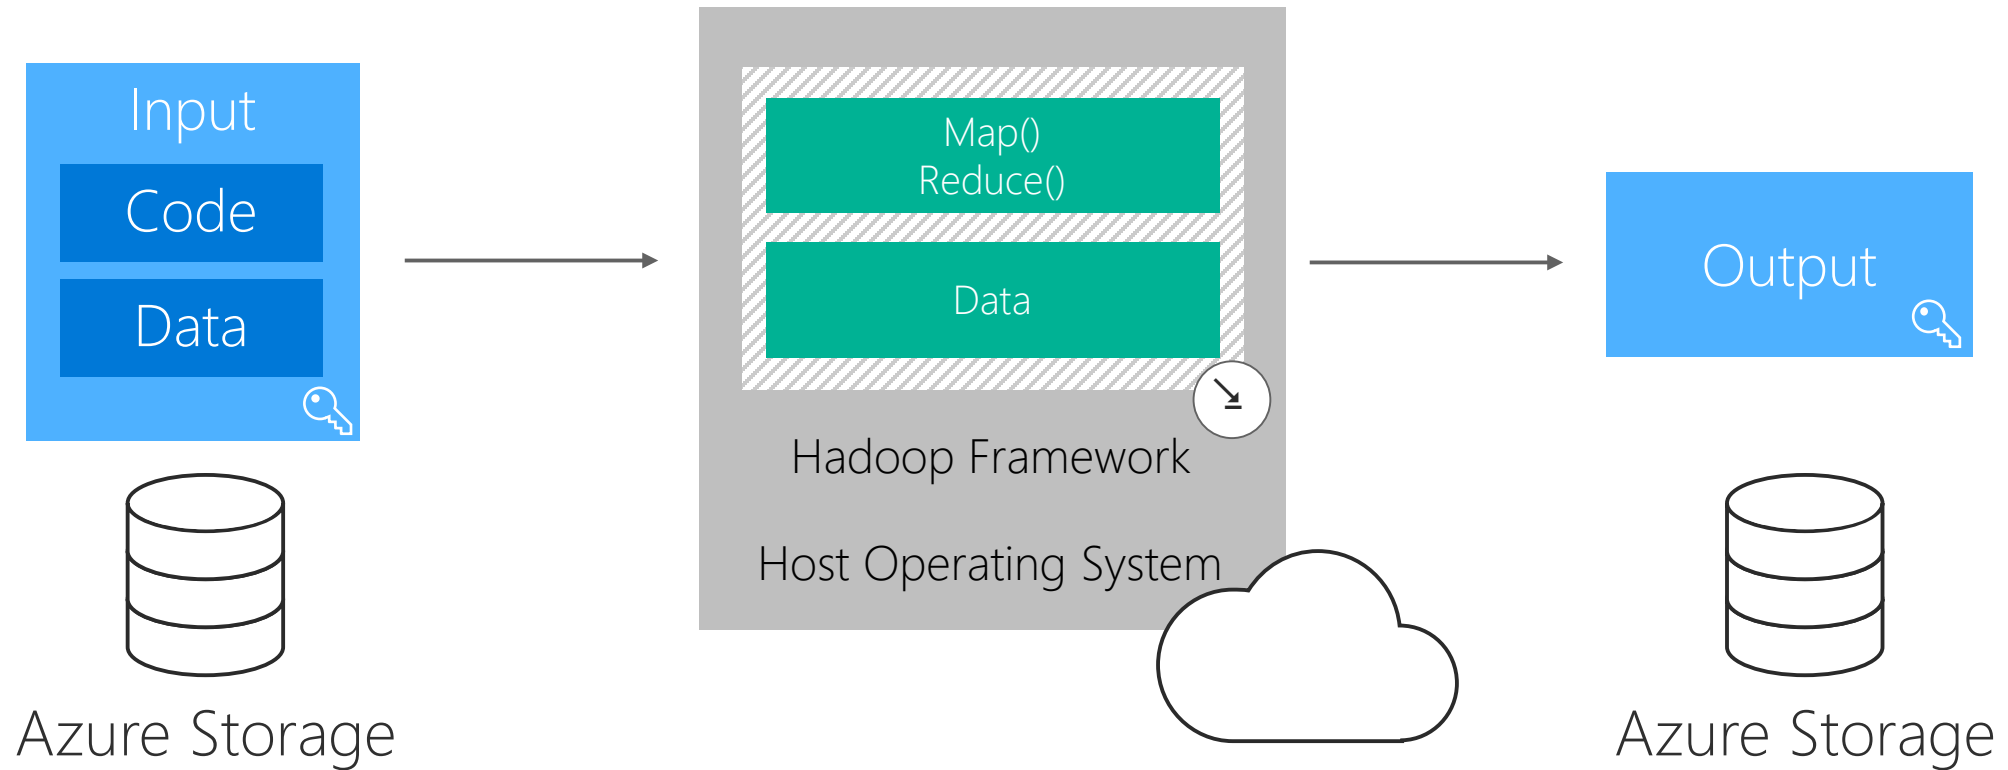
# Demo:

## Coco Ethereum versus Ethereum

# Azure HDInsight



Process massive amounts of data using open source frameworks such as Hadoop, Spark, Hive, Kafka, etc.





# Trusted/Untrusted Code Refactoring for Hadoop



Secret analytics code (C++)

```
void Mapper::map(string k, string v)
{
    /* ... */
}

void Reducer::reduce(
    string k, vector<string> v)
{
    /* ... */
}
```

+

In-Enclave Lib  
(3,300 LLOC)

**TCB = mapred.dll  
+ SGX processor**

**mapred.dll**

Private enclave code



**kcode**

Public enclave code

writeKV() / readKV()

framework.exe

framework.sys

Inside enclave/Trusted

Secret user code

Public generic code

Protocols

5,500 LLOC

Outside enclave/Untrusted

Create enclave

Talk to OS

Bind to Hadoop

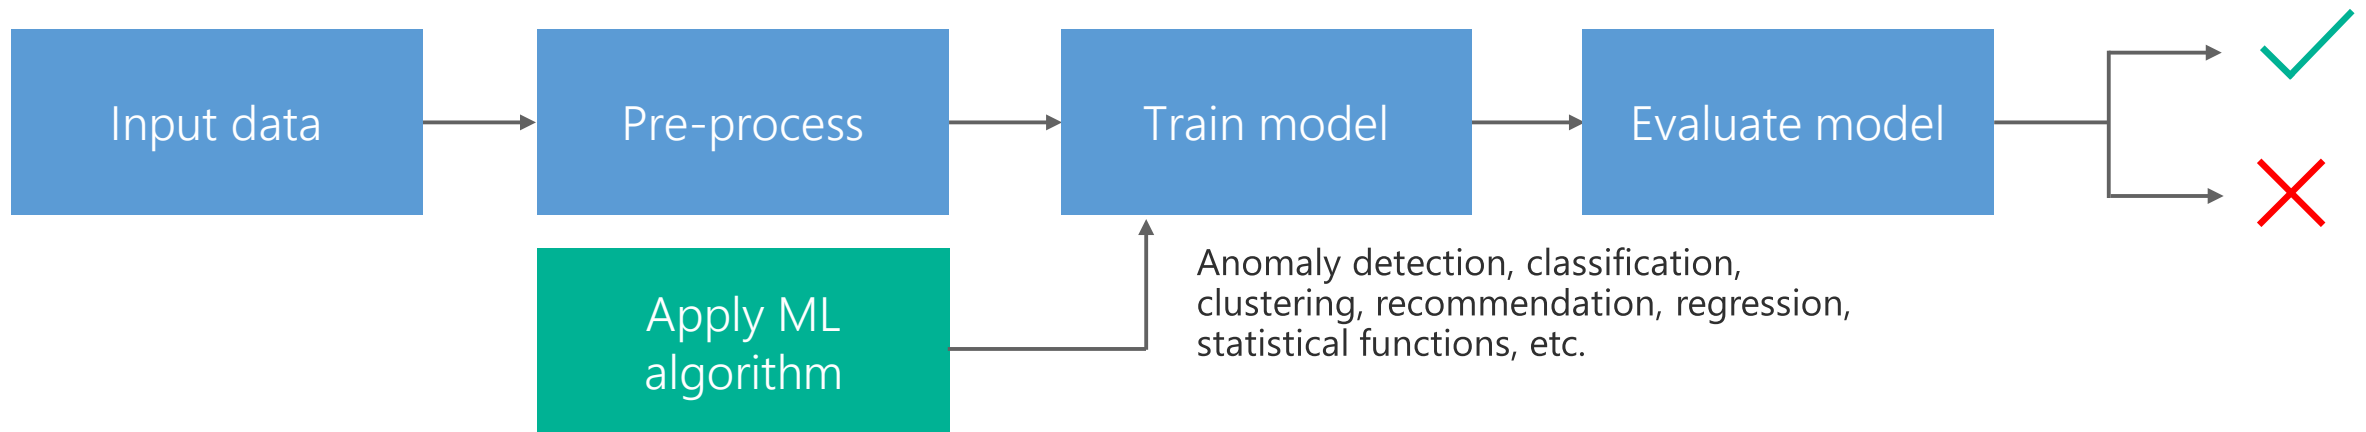
# Azure Machine Learning (ML)



**Machine Learning Studio:** UI web designer to create an end-to-end ML workflow

Large library of predictive analytic algorithms from which to train models

Modules to support data input, output, pre-processing, and visualizations



**Machine Learning API service:** service to deploy prepared ML models at cloud scale and availability

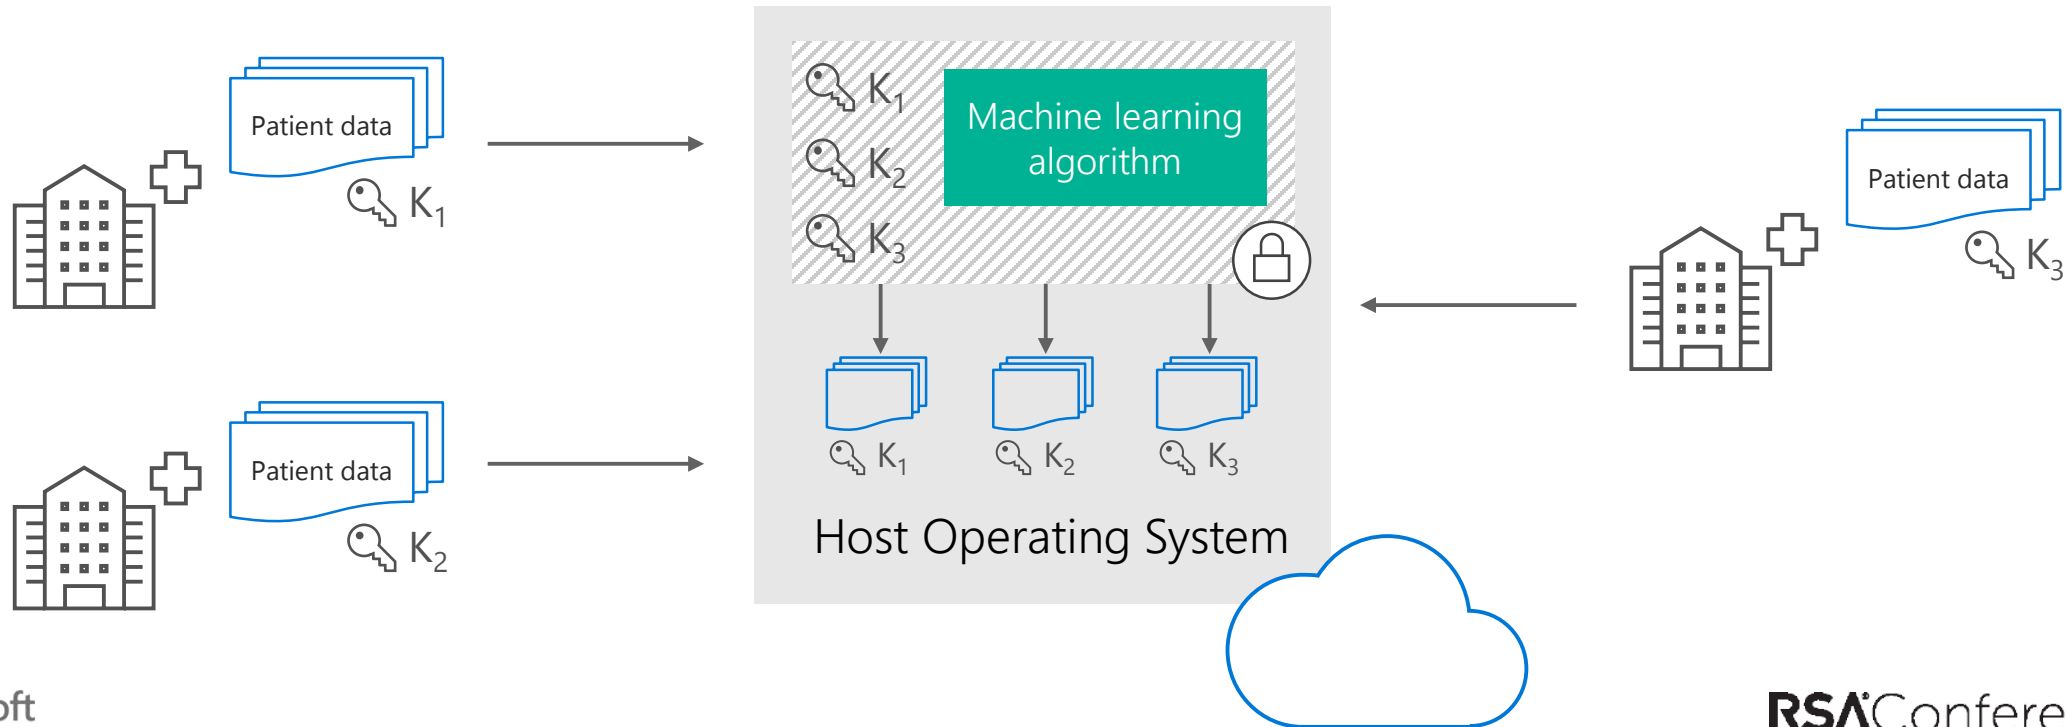
# Confidential multi-party machine learning



Partnered health facilities contribute private patient health data sets to train a ML model

Each facility only sees their respective data sets (aka no one, not even cloud provider, can see all data or trained model, if necessary)

All facilities benefit from using trained model




# Demo:

## Confidential multi-party ML


# Summary



Confidential computing  
in the cloud is in its  
**early stages**



**Microsoft** is driving the  
direction & adoption of  
newer trusted execution  
environments in the cloud



**Azure** is empowering  
new secure business  
scenarios in the cloud

# References

Blockchain with Coco Fx:

<http://aka.ms/cocopaper>

Multi-party machine learning:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/paper.pdf>

SQL Server with Haven:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/osdi2014-haven.pdf>

Map/reduce with VC3:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/vc3-oakland2015.pdf>

Preventing enclave information leaks:

<https://people.eecs.berkeley.edu/~rsinha/research/pubs/pldi2016.pdf>

Using side-channel page faults to extract JPG images:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2017/06/atc17-final230.pdf>