

RSA®Conference2018

San Francisco | April 16 – 20 | Moscone Center

SESSION ID: HTA-F01

ENHANCE VIRTUALIZATION STACK WITH INTEL CET AND MPX



#RSAC

Xiaoning Li

Chief Security Architect
Alibaba Cloud

Ravi Sahita

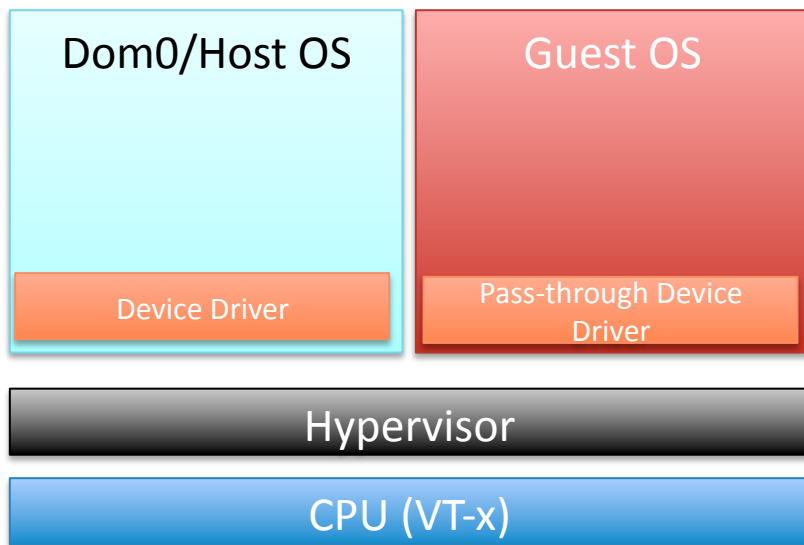
Principal Engineer
Intel Corporation

Agenda

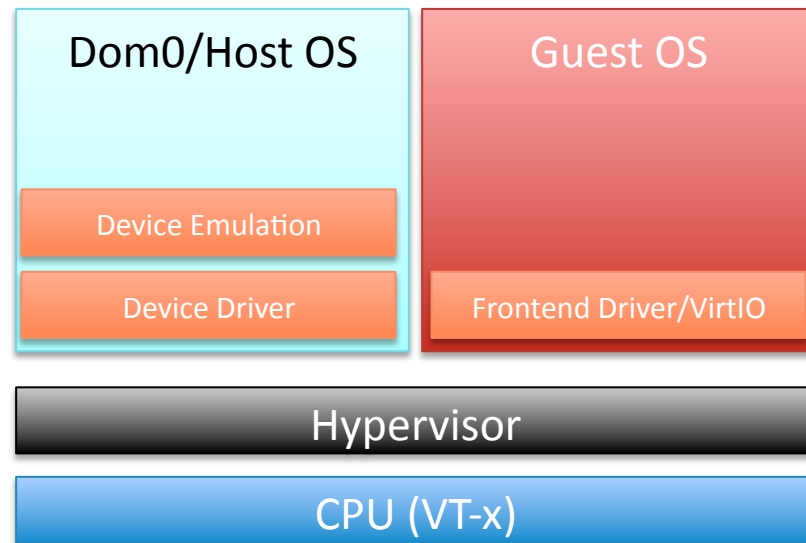


- Full Virtualization Stack
- QEMU Vulnerabilities
- Intel CET
- Intel MPX and PKU
- VM Escape Case Study
- Mitigation with CET/MPX
- Other Mitigations

Full Virtualization Stack



HVM Guest OS with Device Pass-through



HVM Guest OS Without Device Pass-through

Why Device Emulation?




- Supports more guest devices with virtual devices
 - If physical devices number are not enough - limited GPU resource
 - If physical devices don't support device virtualization – Not every device support SRIOV
 - If physical devices don't exist – some outdated devices
- Popular usage in cloud environment to support many VMs
- QEMU can provide device emulation for KVM/XEN, but it brings new attack surface on virtualization stack

RSA®Conference2018



QEMU VULNERABILITIES

QEMU Vulnerabilities by 2018 Jan



CVE-2015-5165 – Information Leak



#RSAC

```
if (proto == ETH_P_IP)
{
    DPRINTF("+++ C+ mode has IP packet\n");

    /* not aligned */
    eth_payload_data = saved_buffer + ETH_HLEN;
    eth_payload_len = saved_size - ETH_HLEN;

    ip = (ip_header*)eth_payload_data;

    if (IP_HEADER_VERSION(ip) != IP_HEADER_VERSION_4) {
        DPRINTF("+++ C+ mode packet has bad IP version %d "
            "expected %d\n", IP_HEADER_VERSION(ip),
            IP_HEADER_VERSION_4);
        ip = NULL;
    } else {
        hlen = IP_HEADER_LENGTH(ip);
        ip_protocol = ip->ip_p;
        ip_data_len = be16_to_cpu(ip->ip_len) - hlen;
    }
} « end if proto==ETH_P_IP »
```

```
struct ip_header {
    uint8_t ip_ver_len; /* version and header length */
    uint8_t ip_tos; /* type of service */
    uint16_t ip_len; /* total length */
    uint16_t ip_id; /* identification */
    uint16_t ip_off; /* fragment offset field */
    uint8_t ip_ttl; /* time to live */
    uint8_t ip_p; /* protocol */
    uint16_t ip_sum; /* checksum */
    uint32_t ip_src, ip_dst; /* source and destination address */
};

#define IP_HEADER_LENGTH(ip) (((ip->ip_ver_len)&0xf) << 2)

/* ETH_MTU = ip header len + tcp header len + payload */
int tcp_data_len = ip_data_len - tcp_hlen;
int tcp_chunk_size = ETH_MTU - hlen - tcp_hlen;
```


CVE-2015-7504 – Heap Overflow



```
} else if (s->looptest == PCNET_LOOPTEST_CRC ||  
           ! CSR_DXMTFCS(s) || size < MIN_BUF_SIZE+4) {  
    uint32_t fcs = ~0;  
    uint8_t *p = src;  
  
    while (p != &src[size])  
        CRC(fcs, *p++);  
    *(uint32_t *)p = htonl(fcs);  
    size += 4;  
} else {  
    uint32_t fcs = ~0;  
    uint8_t *p = src;  
  
    while (p != &src[size-4])  
        CRC(fcs, *p++);  
    crc_err = (*(uint32_t *)p != htonl(fcs));  
}
```



RSA®Conference2018



#RSAC

INTEL CET

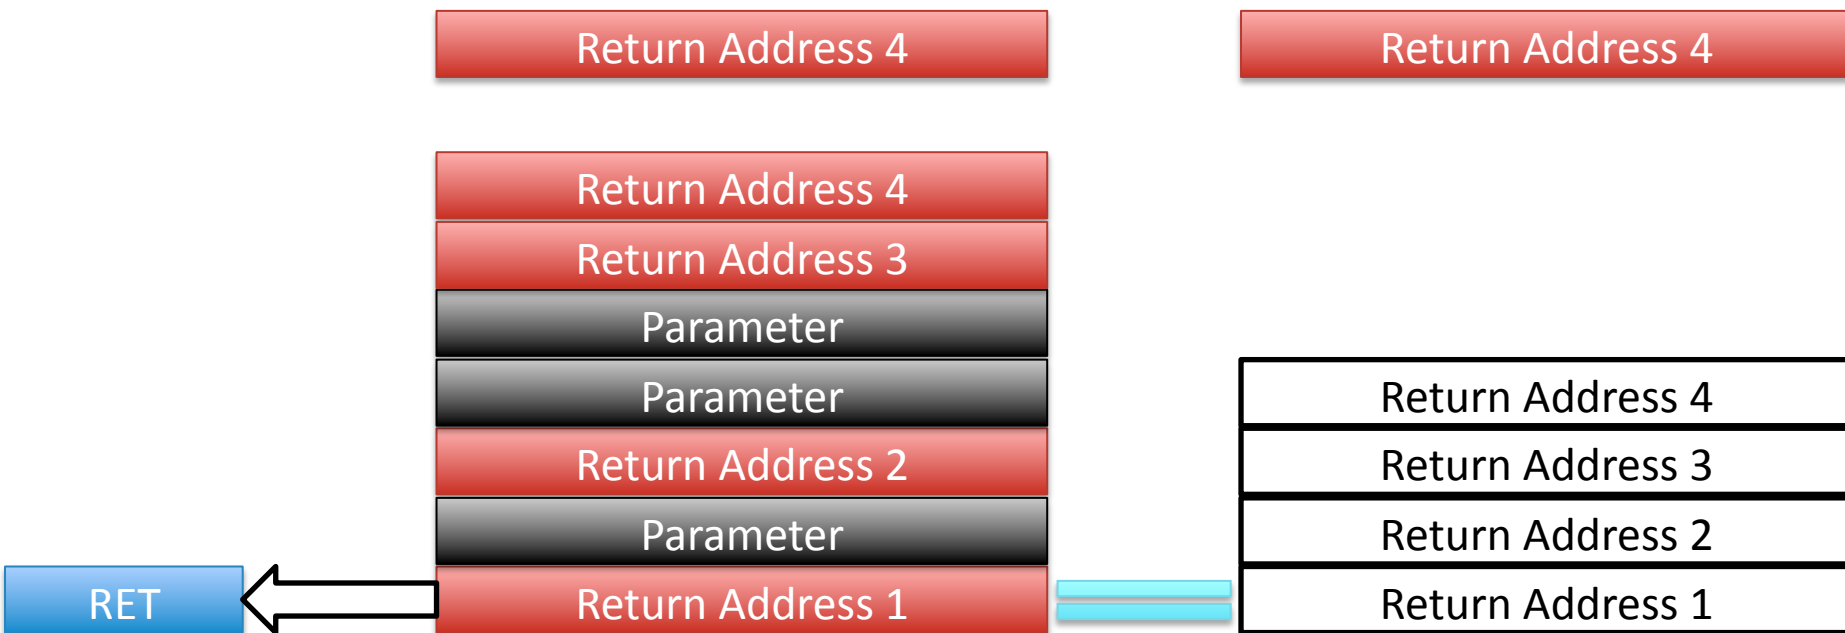
Intel® Control-flow Enforcement Technology



Shadow Stack



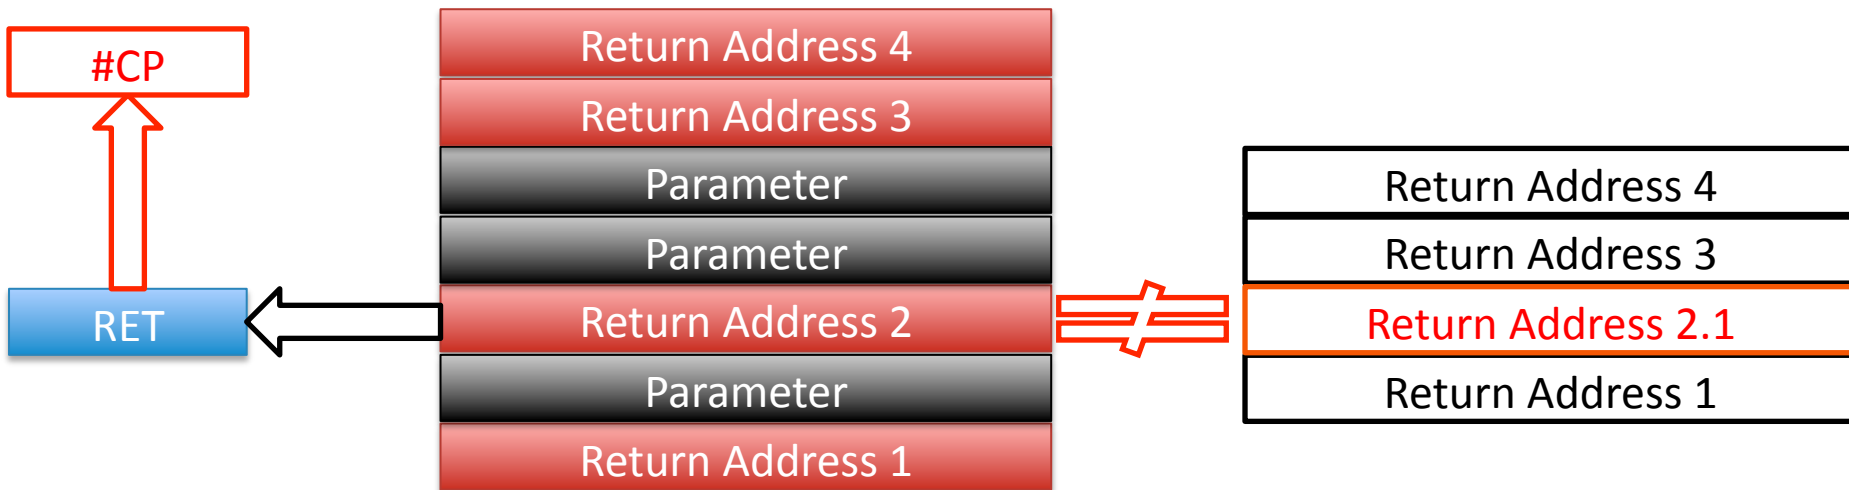
#RSAC



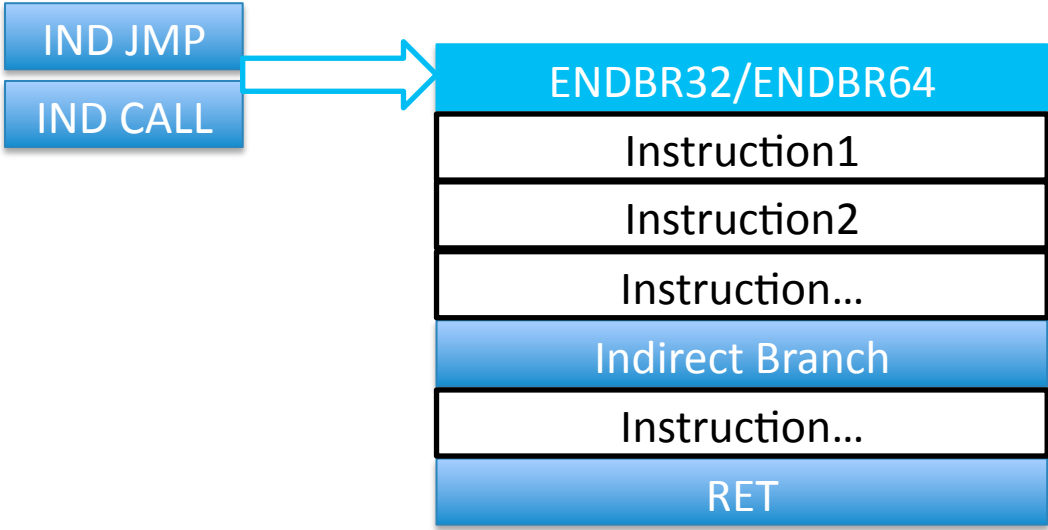
Shadow Stack Control Protection Exception



#RSAC



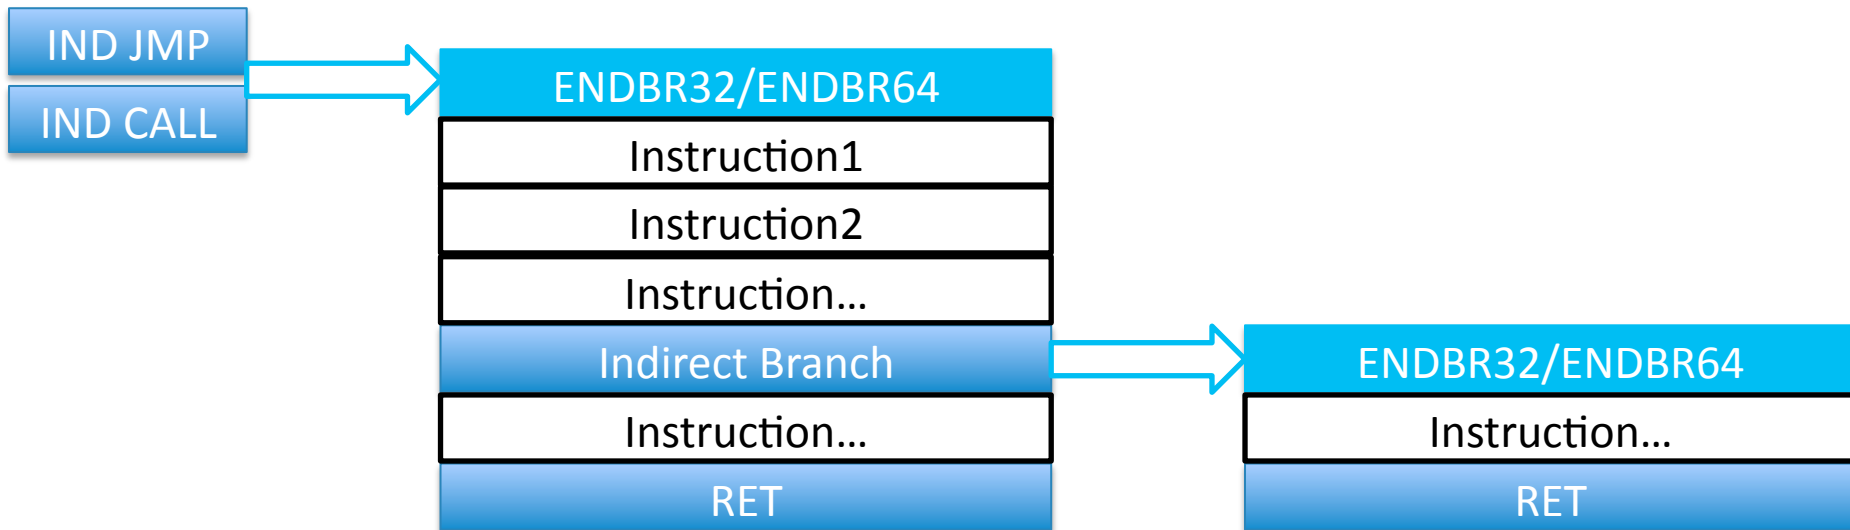
Indirect Branch Tracking



Indirect Branch Tracking



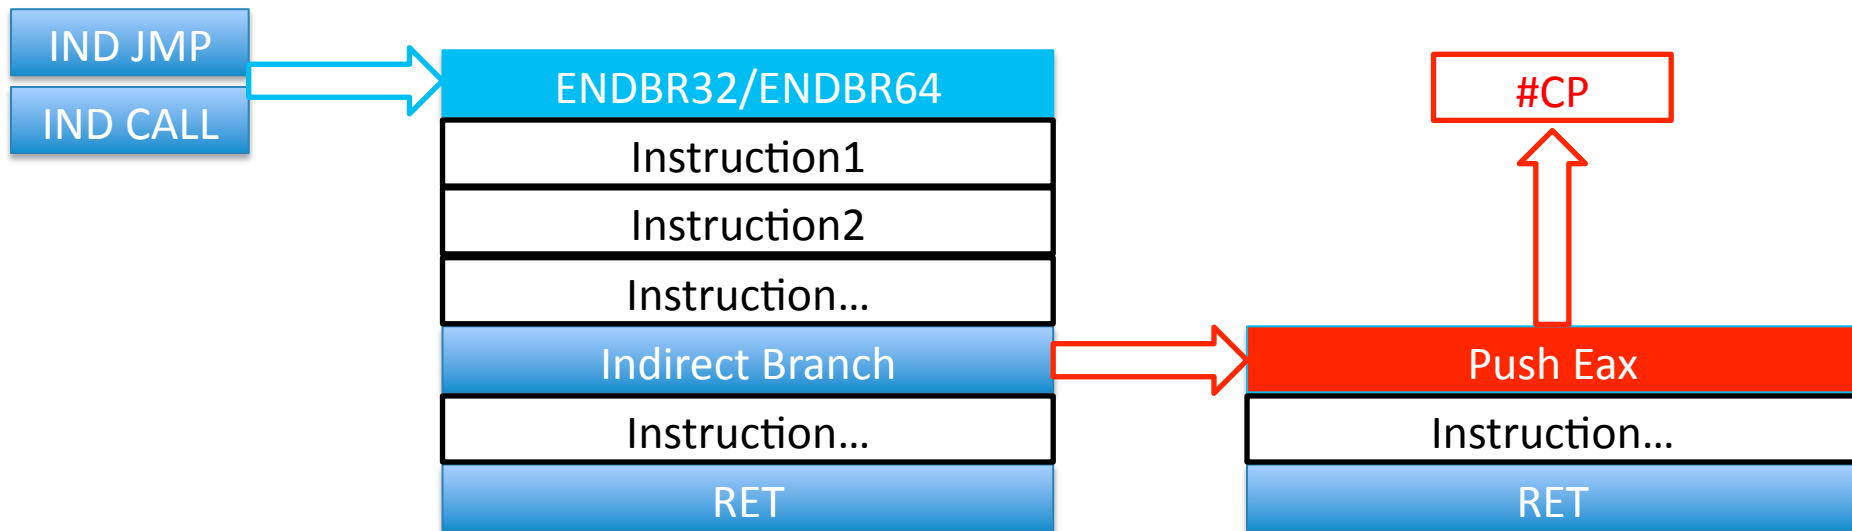
#RSAC



Indirect Branch Tracking #CP



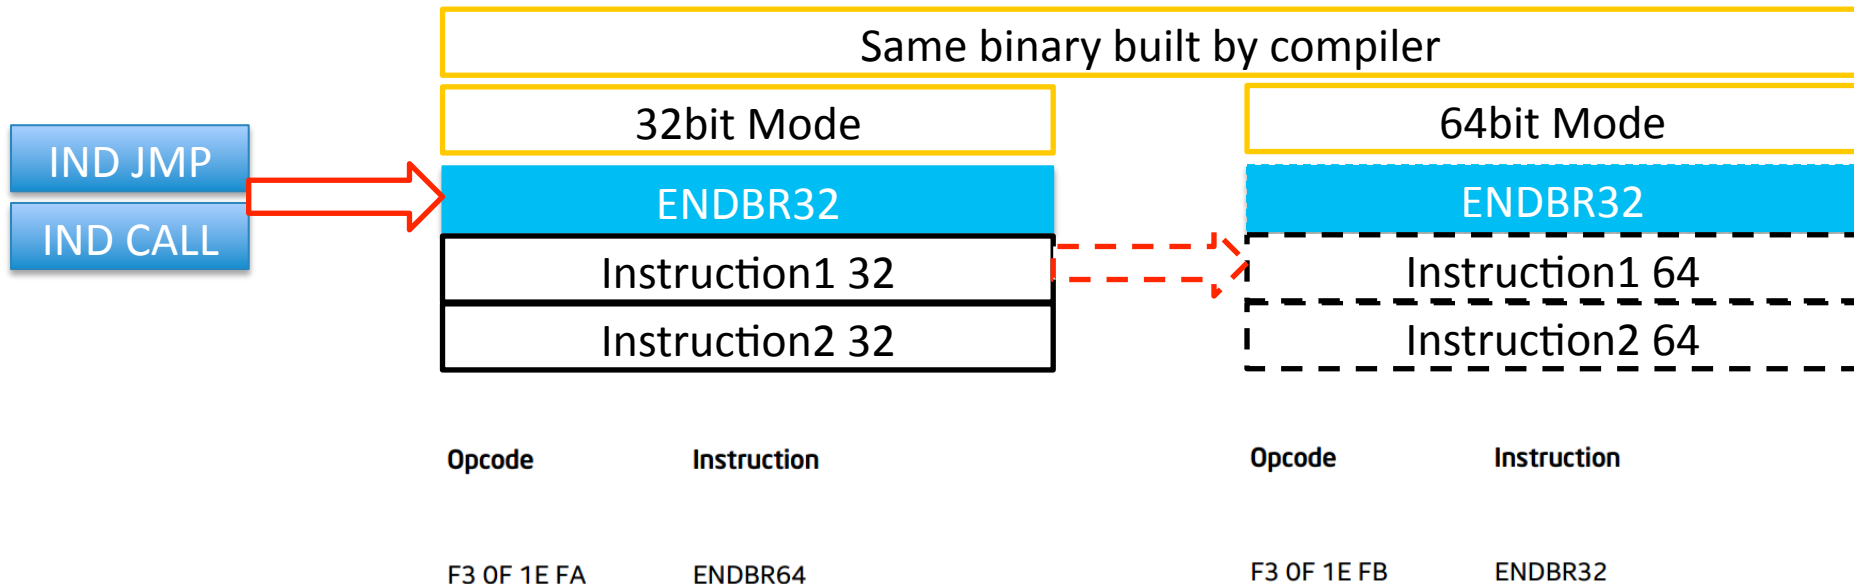
#RSAC



Cross Mode Indirect Branch Tracking



#RSAC




RSA®Conference2018



#RSAC

INTEL MPX

Intel® Memory Protection Extensions



Bound Table

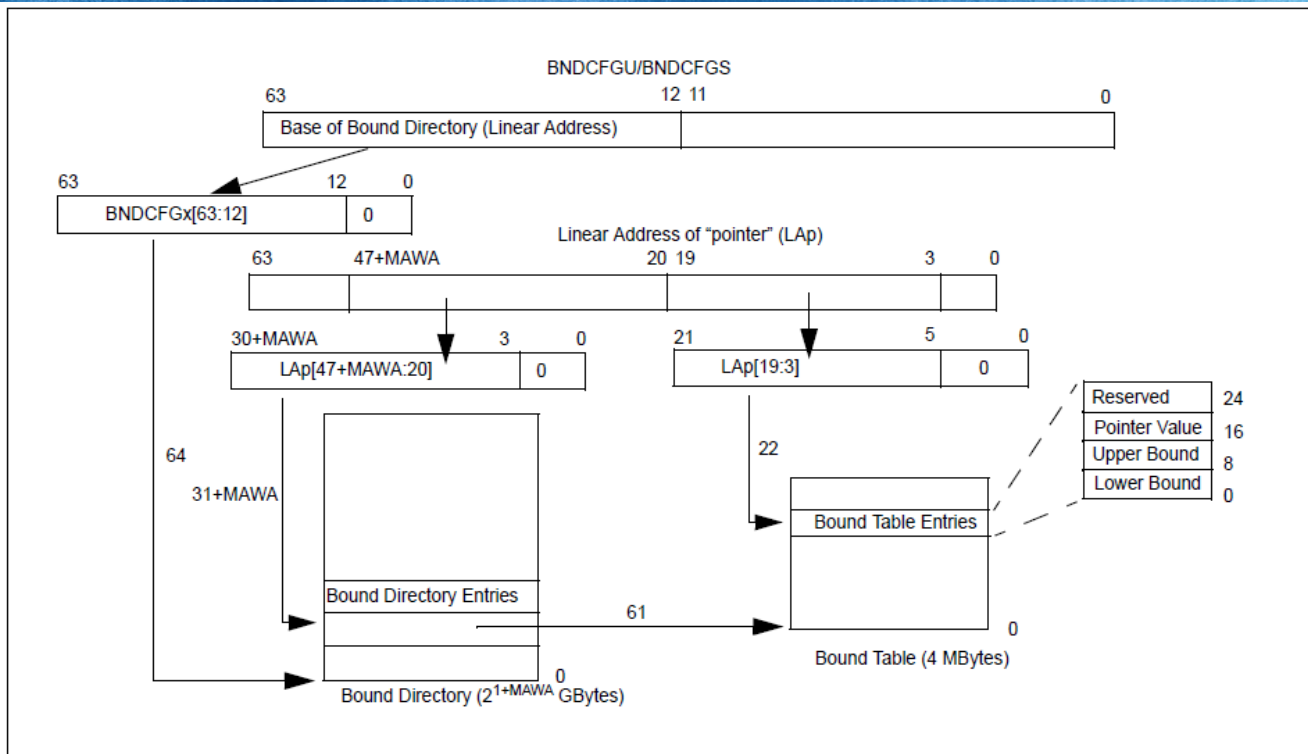


Figure 17-4. Bound Paging Structure and Address Translation in 64-Bit Mode

Bound Instructions



BNDMK	Create a LowerBound and a UpperBound in a register.
BNDCL	Check the address of a memory reference against a LowerBound.
BNDCU	Check the address of a memory reference against an UpperBound in 1's compliment form.
BNDCN	Check the address of a memory reference against an UpperBound not in 1's compliment form.
BNDMOV	Copy or load from memory of the LowerBound and UpperBound to a register.
BNDMOV	Store to memory of the LowerBound and UpperBound from a register.
BNDLDX	Load bounds using address translation.
BNDSTX	Store bounds using address translation.

RSA®Conference2018



#RSAC

QEMU VM ESCAPE CASE STUDY

CVE-2015-5165 and CVE-2015-7504 Exploit



Escape From The Docker-KVM-QEMU Machine

Shengping Wang, Xu Liu
Qihoo 360 Marvel Team

#HITB2016AMS D1T1 - Escape From The Docker KVM QEMU Machine - Shengping Wang and Xu Liu

Title : VM escape - QEMU Case Study

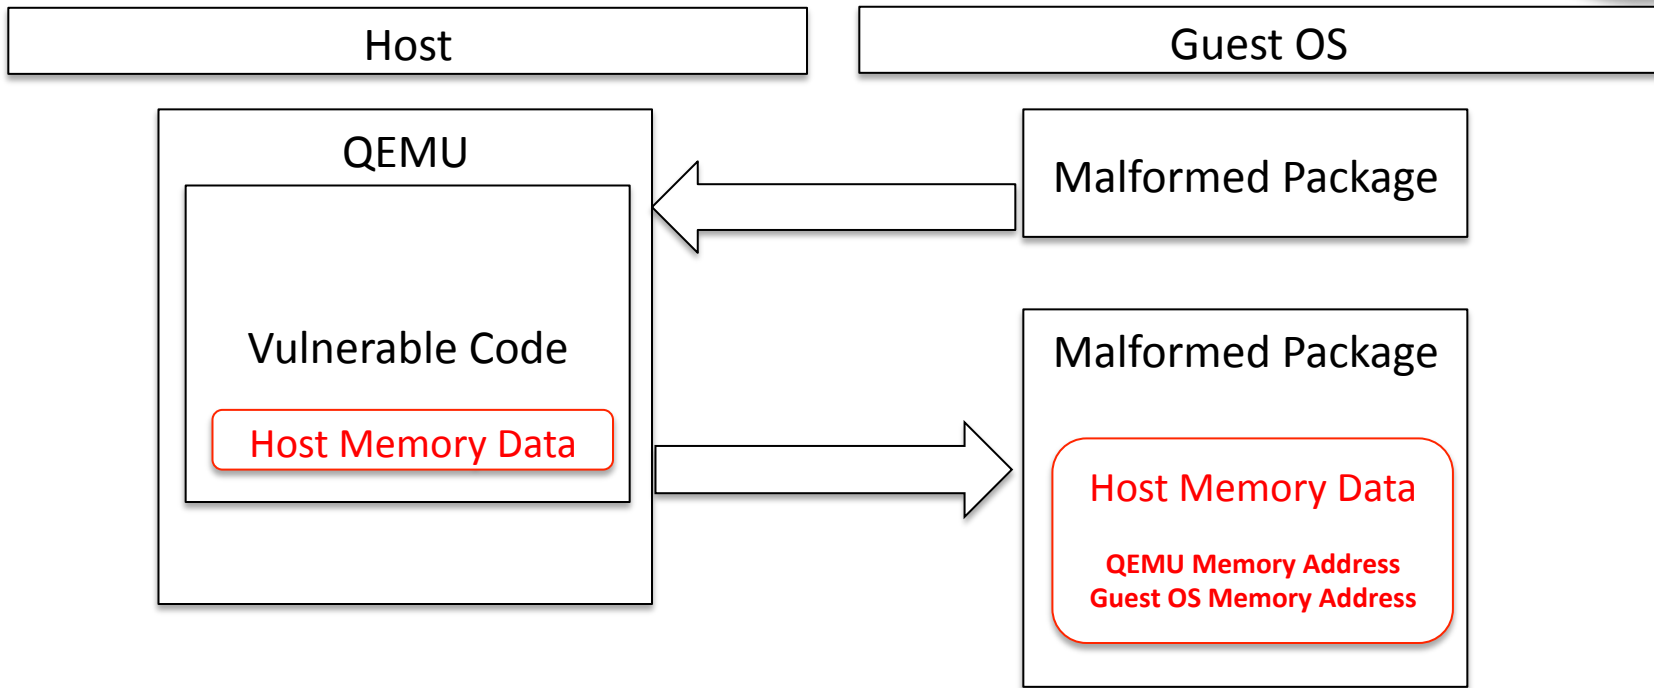
Author : Mehdi Talbi & Paul Fariello

Date : April 28, 2017

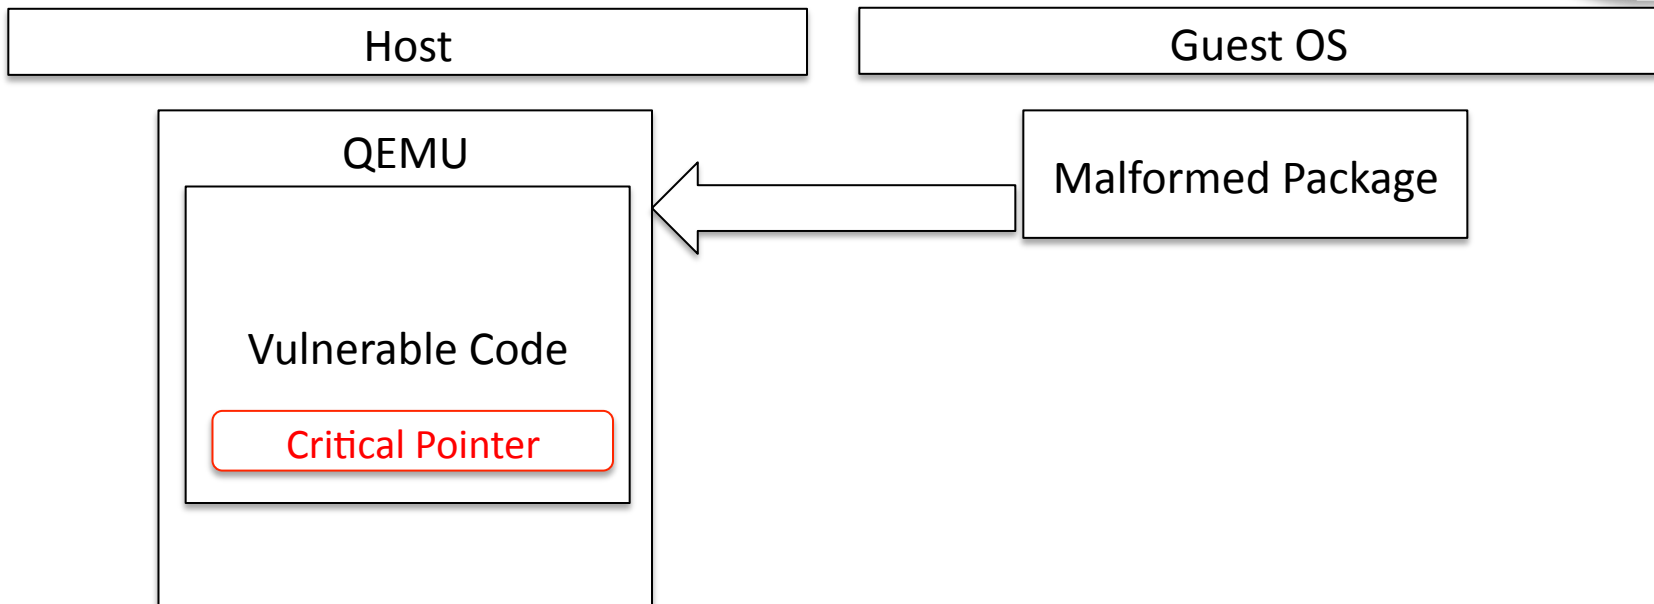
```
|=====|  
|-----=[ VM escape ]-----|  
|=====|  
|-----=[ QEMU Case Study ]-----|  
|=====|  
|-----=[ Mehdi Talbi ]-----|  
|-----=[ Paul Fariello ]-----|  
|=====|
```



CVE-2015-5165 Memory Disclosure



CVE-2015-7504 Code Execution



CVE-2015-7504 Code Execution



#RSAC

```
struct PCNetState_st {
    NICState *nic;
    NICConf conf;
    QEMUTimer *poll_timer;
    int rap, isr, lnkst;
    uint32_t rdra, tdra;
    uint8_t prom[16];
    uint16_t csr[128];
    uint16_t bcr[32];
    int xmit_pos;
    uint64_t timer;
    MemoryRegion mmio;
    uint8_t buffer[4096];
    qemu_irq irq;
    void (*phys_mem_read)(void *dma_opaque, hwaddr addr,
        uint8_t *buf, int len, int do_bswap);
    void (*phys_mem_write)(void *dma_opaque, hwaddr addr,
        uint8_t *buf, int len, int do_bswap);
    void *dma_opaque;
    int tx_busy;
    int looptest;
} « end PCNetState_st » ;
```

```
struct IRQState {
    Object parent_obj;

    qemu_irq_handler handler;
    void *opaque;
    int n;
};
```

```
void qemu_set_irq(qemu_irq irq, int level)
{
    if (! irq)
        return;

    irq->handler(irq->opaque, irq->n, level);
}
```

CVE-2015-7504 Code Execution



```
void qemu_set_irq(qemu_irq irq, int level)
{
    if (!irq)
        return;

    irq->handler(irq->opaque, irq->n, level);
}
```

CVE-2015-7504 Exploit in HITB



```
void qemu_set_irq(qemu_irq irq, int level)
{
    if (!irq)
        return;

    irq->handler(irq->opaque, irq->n, level);
}
```

```
xchg rax, rsp;
ret
```

CVE-2015-7504 Exploit in Phrack



```
void qemu_set_irq(qemu_irq irq, int level)
{
    if (!irq)
        return;

    irq->handler(irq->opaque, irq->n, level);
}
```

qemu_set_irq
Mprotect

shellcode

RSA®Conference2018



#RSAC

MITIGATION WITH CET/MPX

CVE-2015-5165 – Out of Bound Access Defense with MPX



```
for (tcp_send_offset = 0; tcp_send_offset < tcp_data_len; tcp_send_offset += tcp_chunk_size)
{
    uint16_t chunk_size = tcp_chunk_size;

    /* check if this is the last frame */
    if (tcp_send_offset + tcp_chunk_size >= tcp_data_len)
    {
        is_last_frame = 1;
        chunk_size = tcp_data_len - tcp_send_offset;
    }

    DPRINTF("+++ C+ mode TSO TCP seqno %08x\n",
            be32_to_cpu(p_tcp_hdr->th_seq));

    /* add 4 TCP pseudoheader fields */
    /* copy IP source and destination fields */
    memcpy(data_to_checksum, saved_ip_header + 12, 8);

    DPRINTF("+++ C+ mode TSO calculating TCP checksum for "
            "packet with %d bytes data\n", tcp_hlen +
            chunk_size);

    if (tcp_send_offset)
    {
        memcpy((uint8_t*)p_tcp_hdr + tcp_hlen, (uint8_t*)p_tcp_hdr + tcp_hlen + tcp_send_offset, chunk_size);
    }
}
```

Enable MPX on packet memory access

```
struct PCNetState_st {
    NICState *nic;
    NICConf conf;
    QEMUTimer *poll_timer;
    int rap, isr, lnkst;
    uint32_t rdra, tdra;
    uint8_t prom[16];
    uint16_t csr[128];
    uint16_t bcr[32];
    int xmit_pos;
    uint64_t timer;
    MemoryRegion mmio;
    uint8_t buffer[4096];
    qemu_irq irq;
}
```

RSA Conference 2018

CVE-2015-7504 Exploit Defense with CET



```
void qemu_set_irq(qemu_irq irq, int level)
{
    if (!irq)
        return;

    irq->handler(irq->opaque, irq->n, level);
}
```

```
xchg rax, rsp;
ret
```

Shadow Stack can stop “xchg rax, rsp; ret”

CVE-2015-7504 Exploit Defense with CET



```
void qemu_set_irq(qemu_irq irq, int level)
{
    if (!irq)
        return;

    irq->handler(irq->opaque, irq->n, level);
}
```

qemu_set_irq
Mprotect

shellcode

Indirect Branch Tracking can stop `irq->handler` calling `qemu_set_irq` without valid tag

RSA®Conference2018



#RSAC

OTHER MITIGATIONS

Other Mitigation - PKU



62:59	Protection key; if CR4.PKE = 1, determines the protection key of the page (see Section 4.6.2); ignored otherwise																													
-------	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	Bit Position
W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	W	A	
D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	D	
15	15	14	14	13	13	12	12	11	11	10	10	9	9	8	8	7	7	6	6	5	5	4	4	3	3	2	2	1	1	0	0	

Figure 2-9. Protection Key Rights Register for User Pages (PKRU)

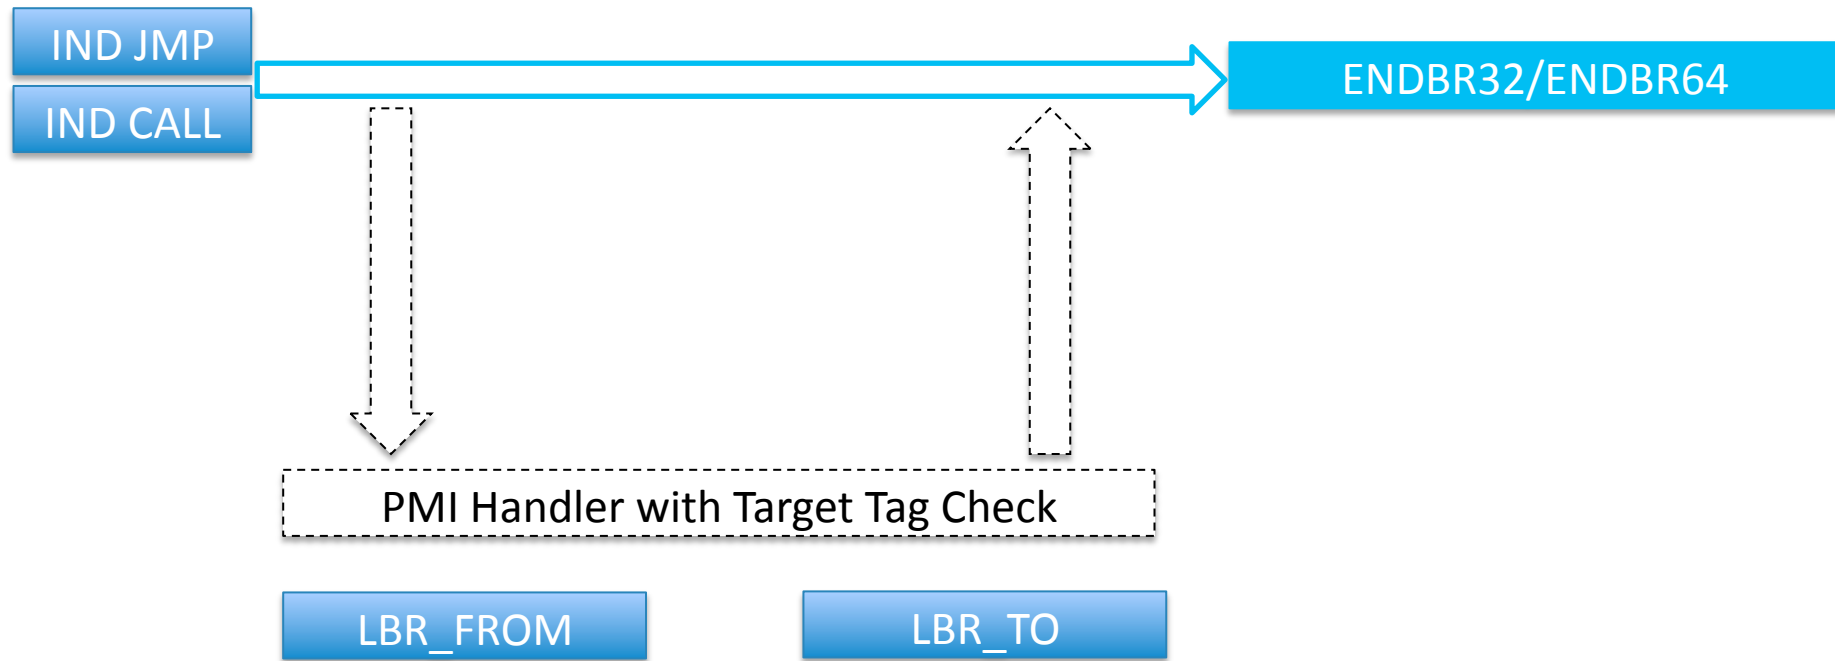
RDPKRU—Read Protection Key Rights for User Pages

Opcode*	Instruction	Op/En	64/32bit Mode Support	CPUID Feature Flag
NP OF 01 EE	RDPKRU	Z0	V/V	OSPKE

WRPKRU—Write Data to User Page Key Register

Opcode*	Instruction	Op/En	64/32bit Mode Support	CPUID Feature Flag
NP OF 01 EF	WRPKRU	Z0	V/V	OSPKE

Other Mitigation - PMI



Summary



- VM escape is practical and impacts cloud foundation security
- Intel CET/MPX can enhance mitigation on ROP/JOP/COP and buffer overflow, specifically on cloud virtualization stack

Call For Actions



- Apply new CPU mechanisms such as CET/MPX/PMU/PKU on exploit defense

RSA®Conference2018



#RSAC

THANK YOU! QUESTIONS?

XIAONING.LI@ALIBABA-INC.COM