

**RSA**®Conference2018

San Francisco | April 16 – 20 | Moscone Center

SESSION ID: PDAC-T07

# SPECTRE ATTACKS: EXPLOITING SPECULATIVE EXECUTION

**Paul Kocher**

[www.paulkocher.com](http://www.paulkocher.com)



#RSAC



## Meltdown



Exposes physical memory

Malicious user process

None

Intel, a few ARM

Complete fix via O/S update

## Spectre



Exposes memory from kernel, other processes, or sandbox


Malicious user process or sandboxed code (JavaScript...)

Tailor to victim SW (comparable to buffer overflow)

Virtually all fast CPUs – Intel, ARM, AMD, Power, Sparc...

Messy partial fixes, if any  
x86: Microcode + O/S + drivers + application  
ARM: varies/none

# Squeezing out CPU performance



Single-thread speed gains require getting more done per clock cycle

- Memory latency is slow and not improving much
- Clock rates are maxed out: Pentium 4 was 3.8 GHz in 2004

How to do more per clock cycle?

- Reducing memory delays → Caches
- Working during delays → Speculative execution

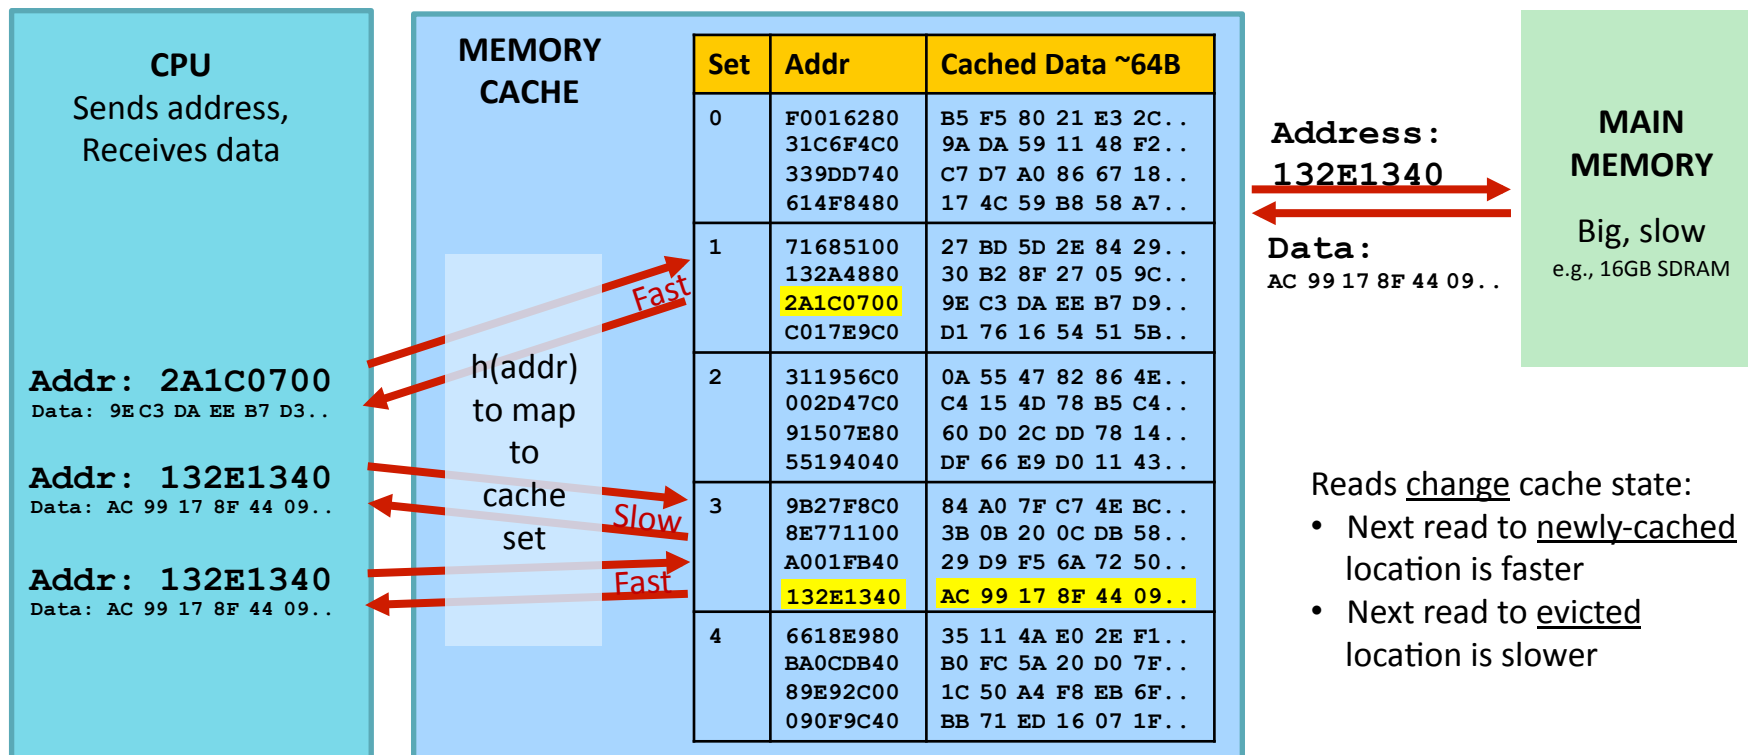
} Spectre leverages interplay  
of these optimizations



# Memory caches



- Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



# Speculative execution



CPU's run instructions

Correct result is defined as the result of performing instructions in-order

CPU's may run instructions out-of-order if this doesn't affect result

Example:

```
a ← constant  
b ← slow_to_obtain  
c ← f(a)  // start before b finishes
```

CPU's can also *guess* likely program path and do speculative execution

Example:

```
if (uncached_value_usually_1 == 1)  
    compute_something()
```

- ▶ Branch predictor guesses that if() will be 'true' based on prior history (not current operation)
- ▶ Starts executing compute\_something() speculatively -- but doesn't save changes
- ▶ When if() can be evaluated definitively, check if guess was correct:
  - ▶ Correct: Save speculative work – performance gain
  - ▶ Incorrect: Discard speculative work

# Architectural World

Guarantees final register values and memory state

Software security assumes CPU correctness.  
Do speculation errors violate this assumption?


Set up the conditions so the processor will make a desired mistake

Fetch the sensitive data from the covert channel

# Speculative World

CPU regularly performs incorrect calculations, but deletes its mistakes

Mistake leaks sensitive data into a covert channel (e.g., state of the cache)



# Conditional branch (Variant 1) attack



```
if (x < array1_size)
    y = array2[array1[x]*512];
```

Imagine this code was in a kernel API where `x` came from caller (e.g., attacker)

Execution without speculation is safe:

CPU will not evaluate `array2[array1[x]*512]` unless `x < array1_size`

What about speculative execution?

# Conditional branch (Variant 1) attack



```
if (x < array1_size)
    y = array2[array1[x]*512];
```

## Attacker objective

- Read sensitive memory
- Example reads `array1[N+8]`  
... where `N+8 > array1_size`

## Attack setup:

- Train branch predictor to expect `if()` is true  
(e.g., call with `x < array1_size`)
- Evict `array1_size` and `array2[]` from cache

## Memory & Cache Status

`array1_size = 00000008`

Memory following `array1` base address:

8 bytes of data (value doesn't matter)

[... N bytes of other memory ...]

**09** F1 98 CC 90 ... (something secret)

`array2[ 0*512]`  
`array2[ 1*512]`  
`array2[ 2*512]`  
`array2[ 3*512]`  
`array2[ 4*512]`  
`array2[ 5*512]`  
`array2[ 6*512]`  
`array2[ 7*512]`  
`array2[ 8*512]`  
`array2[ 9*512]`  
`array2[10*512]`  
`array2[11*512]`  
...

Contents don't matter  
only care about cache **status**

Uncached

Cached



# Conditional branch (Variant 1) attack



```
if (x < array1_size)
    y = array2[array1[x]*512];
```

Attacker calls victim code with  $x=(N+8)$

- Speculative exec while waiting for `array1_size`
  - Predict that `if()` is true
  - Read address (`array1` base +  $x$ ) w/ out-of-bounds  $x$
  - Read returns secret byte = **09** [fast – in cache]
  - Request memory at (`array2` base +  $09*512$ )

## Memory & Cache Status

`array1_size` = 00000008

Memory following `array1` base address:  
8 bytes of data (value doesn't matter)  
[... N bytes of other memory ...]

**09 F1 98 CC 90** ... (something secret)

```
array2[ 0*512]
array2[ 1*512]
array2[ 2*512]
array2[ 3*512]
array2[ 4*512]
array2[ 5*512]
array2[ 6*512]
array2[ 7*512]
array2[ 8*512]
array2[ 9*512]
array2[10*512]
array2[11*512]
...
```

Contents don't matter  
only care about cache **status**

Uncached

Cached

# Conditional branch (Variant 1) attack



```
if (x < array1_size)
    y = array2[array1[x]*512];
```

Attacker calls victim code with  $x=(N+8)$

- Speculative exec while waiting for `array1_size`
  - Predict that `if()` is true
  - Read address (`array1` base +  $x$ ) w/ out-of-bounds  $x$
  - Read returns secret byte = **09** [fast – in cache]
  - Request memory at (`array2` base +  $09*512$ )
  - Brings `array2[09*512]` into the cache
  - Realize `if()` is false: discard speculative work
- Finish operation & return to caller

Attacker times reads from `array2[i*512]`

- Read for  $i=09$  is fast (cached), revealing secret byte

## Memory & Cache Status

`array1_size` = 00000008

Memory following `array1` base address:  
8 bytes of data (value doesn't matter)  
[... N bytes of other memory ...]

**09 F1 98 CC 90** ... (something secret)

```
array2[ 0*512]
array2[ 1*512]
array2[ 2*512]
array2[ 3*512]
array2[ 4*512]
array2[ 5*512]
array2[ 6*512]
array2[ 7*512]
array2[ 8*512]
array2[ 9*512]
array2[10*512]
array2[11*512]
...
```

Contents don't matter  
only care about cache **status**

Uncached

Cached

# Violating JavaScript's sandbox



- JavaScript code runs in a sandbox
  - Not permitted to read arbitrary memory
  - No pointers, array accesses are bounds checked
- Browser runs JavaScript from untrusted websites
  - JavaScript engine can interpret code (slow) or compile it (JIT) to run faster
  - In all cases, engine must enforce sandbox (e.g., apply bounds checks)
- Speculative execution can blast through safety checks...
  - Can we write JavaScript that compiles into machine code that leaks memory contents?

# Violating JavaScript's sandbox



`index` will be in-bounds on training passes, and out-of-bounds on attack passes

JIT knows this check ensures `index < length`, so it omits bounds check in next line. Caller evicts `length` for attack passes.

```
if (index < simpleByteArray.length) {  
  index = simpleByteArray[index | 0];  
  index = (((index * TABLE1_STRIDE) | 0) & (TABLE1_BYTES - 1)) | 0;  
  localJunk ^= probeTable[index | 0] | 0;  
}
```

Do the out-of-bounds read on attack passes!

Hint to JS that result is an integer

4096 bytes (= page size)

Keeps the JIT from adding unwanted bounds checks on the next line

Use the result so computation isn't optimized away

Leak out-of-bounds read result into cache state!

- Branch predictor matches jump history, so bit operations used to load `index` for training vs. attack passes
- No access to `clflush`, so eviction done by bulk reads on 4KB intervals
- JavaScript timers degraded, so timing done via a thread ("worker") that decrements a shared memory location in a tight loop




# Indirect branches




Conditional branches have 2 possible destinations

- E.g., next instruction if false, jump destination if true



- Indirect branches can go anywhere

- Examples on x86: `jmp [1234567]`      `jmp eax`      `ret`
- Branch target buffer: CPU tracks past jump destinations to make quick guesses
- If destination is delayed, CPU guesses and proceeds speculatively



- Vastly more freedom for attacker – billions of possible destinations (vs 2)

# Poisoning indirect branches



- Pick indirect branch to redirect speculatively
  - e.g., a jump that occurs with attacker-controlled values in some registers
- Pick destination for victim to speculatively execute (the “gadget”)
  - e.g., gadget that leaks attacker-chosen memory address into covert channel
  - Like return oriented programming, but gadget also doesn’t have to return nicely
- Attack
  - **Mistrain** branch prediction/BTB so speculative execution will go to gadget
  - **Evict** or flush destination address from the cache to ensure long duration speculative execution
  - **Execute** victim so it runs gadget speculatively
  - **Detect** change in cache state (e.g., EVICT+RELOAD) to determine memory data
  - **Repeat** for more bytes

# Indirect branch attack example



## Victim program

- Example code repeatedly hashes a secret key & header of a file
- Happens to calls Sleep() with attacker-controlled data in **ebx** & **edi**
- First instruction a DLL call (incl. Sleep) is an indirect jump: `jmp dword ptr ds:[754F009Ch]`

## Gadgets

- Searched a simple program's system DLL code to identify gadgets leveraging in **ebx** & **edi**
- Gadget in ntdll.dll:

```
adc    edi,dword ptr [ebx+edx+13BE13BDh]
adc    dl,byte ptr [edi]
```

First instruction adds word from attacker-controllable address [edx=3 in victim, carry is clear] onto edi

Second instruction reads memory location [edi]

- Used an another gadget to break ASLR:

```
sbb    eax,dword ptr [esp+ebx]
```

  - Repeat: Try a value for ebx then check cache state of a chunk of shared readable memory
  - Reveals value of esp

## Attack

- Mistrain jmp to gadget, evict address with jmp destination, execute victim, detect cache state

# Poisoning indirect branches: Notes



- Predictor works on virtual addresses
  - Training using attack process virtual addresses gets applied to victim process
  - Collisions in history tracking allow various simplifications
- Speculative execution only runs from addresses the victim can execute
  - Gadget must be executable in victim's address space
  - Example: gadgets in Windows DLLs
  - Example: gadgets created using eBPF JIT
- Reminder: Spectre leverages victim's permissions
  - Speculative execution occurs in victim context  
[Meltdown is different – privilege escalation]



# Mitigations

Mitigation. noun. “The action of reducing the severity, seriousness, or painfulness of something”

Not necessarily a complete solution



# Mitigations: Conditional branch variant




CPU vendors redefining instructions that work on existing CPUs

- `LFENCE` (Intel, some AMD may need microcode update)
- CSDB for some ARM processors, others have no mitigation other than using conditional loads

Works in theory... what about in practice?

- Must protect all exploitable code patterns: Miss one and attacker can read entire process memory
- Manual insertion is impractical
  - Many codebases have millions of jumps
  - Code reviews are really hard (speculation runs far -- e.g., 188++ instructions)
- Insert everywhere = simple-ish but big performance hit
  - e.g., SHA-256 test = 59.4% performance reduction
  - Static analysis is hard: Choose performance or security
    - Microsoft Visual C/C++ /Qspectre only adds `LFENCE` to a narrow known-bad code pattern
    - 13 of 15 code examples yielded unprotected output <https://www.paulkocher.com/doc/MicrosoftCompilerSpectreMitigation.html>
- Lots of software potentially impacted: Kernel, device drivers, web servers, interpreters/JITs, crypto code, databases...
  - Kernels will get careful attention, but lots other software won't (drivers, etc.)



Allows blaming developer post-breach (“should’ve added an `LFENCE` there...”)

# Mitigations: Indirect branch variant



x86:

- New MSRs created via microcode
  - IBRS: Restrict speculation of indirect branches - e.g., disable predictor
  - STIBP: Prevents hyperthread from training predictor - e.g., stop training
  - IBPB: Prevent prior learning from influencing predictions - e.g., clear predictor state
  - Messy:
    - Microcode updates for many (but not all) Intel & AMD CPUs. Stability problems + retractions...
    - Available from kernel mode only (use by applications unclear)
    - Substantial performance impacts
- Retpoline proposal from Google
  - Replace indirect jumps with construction that resists indirect branch poisoning
  - Works on some CPUs (e.g., Haswell), but need microcode updates for others (e.g., Skylake)

- ARM: No generic mitigation option

- Faster CPUs broadly impacted (e.g., Cortex-A9, A15, A17, A57, A72, A73, A75...)
- On some chips, software may be able to invalidate/disable branch predictor (with “non-trivial performance impact”) <https://developer.arm.com/support/security-update/download-the-whitepaper>

# DOOM with only MOV instructions



## Only MOV instructions

- No branch instructions
- One big loop with exception at the end to trigger restart

## Sub-optimal performance

- One frame every ~7 hours

### A branchless DOOM

This directory provides a branchless, mov-only version of the classic DOOM video game.




*DOOM, running with only mov instructions.*

This is thought to be entirely secure against the Meltdown and Spectre CPU vulnerabilities, which require speculative execution on branch instructions.

<https://github.com/xoreaxeaxe/movfuscator/tree/master/validation/boom>



Because of software bugs, computer security was in a dire situation



Spectre doesn't change the magnitude of the risk, but adds to the mess

Complexity of fixes -> new risks

Psychology of unfixed vulnerabilities

# Is Spectre a bug?



- Everything is working 100% as intended
  - Branch predictor is learning from history, as expected
  - Speculative execution unwinds architectural state correctly
  - Reads are fetching data the victim is allowed to read
  - Caches are allowed to hold state
  - Covert channels known to exist (+ often impractical to eliminate)
- Everything complies with architectural specs from Intel, ARM, etc.

?!

A simple instruction:

```
add rax, [rbx]
```

## Architectural Guarantee

Compute the correct value in rax\*

\* assuming voltage/clock/temp... all in spec

## Microarchitectural Properties

- Computation time
- Change in memory cache state
- Timing impact on other threads  
(e.g., due to use of the memory bus utilization, arithmetic unit...)
- Analog effects (power consumption, heat...)
- . . .

## Architecture ↔ Software security gap

- CPU architecture guarantees are insufficient for security
- Under-specified -> software & HW developers make different assumptions
- No way to know if code is secure on today's chips
- Future chips may be completely different

# Hardware challenges



- Better/faster CPU mitigations
  - Trade-offs between narrow patches (faster) vs. comprehensive fixes (safer)
    - Tricky – lots ways to adjust Spectre attacks to work around fixes
    - Unlikely to abandon speculative execution – too important for performance
  - Roadmap challenge
    - Break trend of increasing use of speculative execution for performance gains?
  - Tempting to punt to software developers... who will fail
    - Complex & virtually untestable – software teams are failing on easier problems
    - Developers can't cope with microarchitectural complexity (LFENCE, cache minutae...)
    - Punt responsibility vs. fix
- Spectre will linger
  - Long device lifetimes
  - Initial countermeasures are narrow
  - New attack techniques tend to evolve (DPA, buffer overflows, use-after-free...)
  - Embedded chips with old ARM designs will keep getting made for a very long time



# Apply: For users/IT departments



- Install updates + more updates + more updates...
  - O/S fix for Meltdown + kernel hardening against Spectre
  - Iterate as new attacks using Spectre arise
- Improve separation
  - Consider segregating sensitive & untrusted workloads on different servers
  - Be especially cautious about Hyperthreading -- expect new side channel attacks
- No viable alternatives in the market (“buy new hardware”)
  - Intel + AMD + ARM + Power + etc. all impacted – none better than the other
  - Vendors touting immunity largely just too small/obsolete/unoptimized to use speculative execution
- Longer-term: Press vendors to deliver practical high-assurance/low-complexity separation
  - Today’s CPUs + hypervisors + compilers all problematic: prioritize performance > security

# Why wasn't Spectre found earlier?



- Attack is obvious... in hindsight:
  - Component behaviors all widely known (e.g., taught in CPU design textbooks)
  - Speculative execution -> queasy feelings even without an exploit
  - Independent discovery seems to be complete coincidence
- Why missed?
  - Dismissal of prior side channel attacks as out of scope (cache, DPA, fault...)
  - Assumption that established architectures are correct
  - Overwhelming complexity narrows focus of separate technical teams
  - Focus on minutiae (bugs)
  - Performance-first culture

## Spectre isn't the first – or last – hardware security vulnerability...

*“The reality is there are probably other things out there like [Spectre] that have been deemed safe for years. Somebody whose  
**mind is sufficiently warped**  
toward thinking about security threats may find other ways to exploit systems which had otherwise been considered completely safe.”*

— Simon Segars (CEO, Arm Holdings)

# HW issues & responsible disclosure



**Notify vendor(s)**



**Embargo while  
vendor fixes bug**



**Public disclosure**

Messy but works... for software

Which vendors? Who decides? Who coordinates?

- › Vendors: logic cores... chips... products... O/S... apps... cloud services...
- › Who else: Governments... big customers... CERT... security experts...

How long? Who decides length?

- › Time to 'fix' may be years (or  $\infty$ )
- › Worst case: Bad guys exploiting & defenders lack info

How?

- › Role of marketing, legal (32+ lawsuits), investor relations, press...?
- › Spectre/Meltdown coordinated embargo (many vendors + researchers)

*"AMD is not susceptible to all three variants. [...] there is a near zero risk to AMD processors at this time."*

I'm 0-for-2 applying to HW issues: **Spectre + Differential Power Analysis (DPA)**  
(DPA affected virtually all smart cards used for payments, pay TV...)

(# people with need-to-know) >> (# who can keep a secret)

- › Both embargos ended abruptly due to leaks
- › Most impacted organizations got no advance warning

**What is the right ethical process for handing hardware vulnerabilities?**



**Today: Costs of insecurity >>> Value of performance gains**



**Technical challenge**

Building foundations that are unlikely to harbor vulnerabilities



**Cultural challenge**

Unlearning 50+ years of obsession with performance

- Current designs are too complex to secure
- Too many constraints: Performance, legacy compatibility...

# A path forward



Today: Same CPU and O/S designs for everything

- Video games, corporate email, wire transfers...
- No tailoring to security vs. performance needs
  - Example: All cores are identical for multi-core x86 CPUs

Need to bifurcate designs into **faster** vs. **safer** flavors

- Can co-exist on the same chip (analogous to ARM's big.LITTLE for power)
- 'Safer' must be much less complex -- not just a different mode like TrustZone/SGX

Urgent need for designs that are unlikely to harbor unknown vulnerabilities

# Q&A



*If the surgery proves unnecessary, we'll  
revert your architectural state at no charge.*