



TENCENT SECURITY CONFERENCE 2018
2018腾讯安全国际技术峰会

A Dive in to Hyper-V Architecture & Vulnerabilities

Nicolas Joly (@n_joly)
Joe Bialek (@JosephBialek)

MSRC Vulnerabilities & Mitigations Team



This presentation is for informational purposes only. MICROSOFT MAKES NO WARRANTIES, EXPRESS OR IMPLIED, IN THIS SUMMARY.

Hyper-V Bug Bounty (as of August 2018)

RCE w/ Exploit
(Guest-to-Host Escape)

\$250,000 (Hypervisor/Kernel)
\$150,000 (User-mode)

RCE
(Guest-to-Host Escape)

\$200,000 (Hypervisor/Kernel)
\$100,000 (User-mode)

Information Disclosure

\$25,000 (Hypervisor/Kernel)
\$15,000 (User-mode)

Denial of Service

\$15,000 (Hypervisor/Kernel)

See aka.ms/bugbounty for details



TENCENT SECURITY CONFERENCE 2018
2018腾讯安全国际技术峰会

Architecture Overview

(From the perspective of a security researcher who wants to find guest to host bugs)

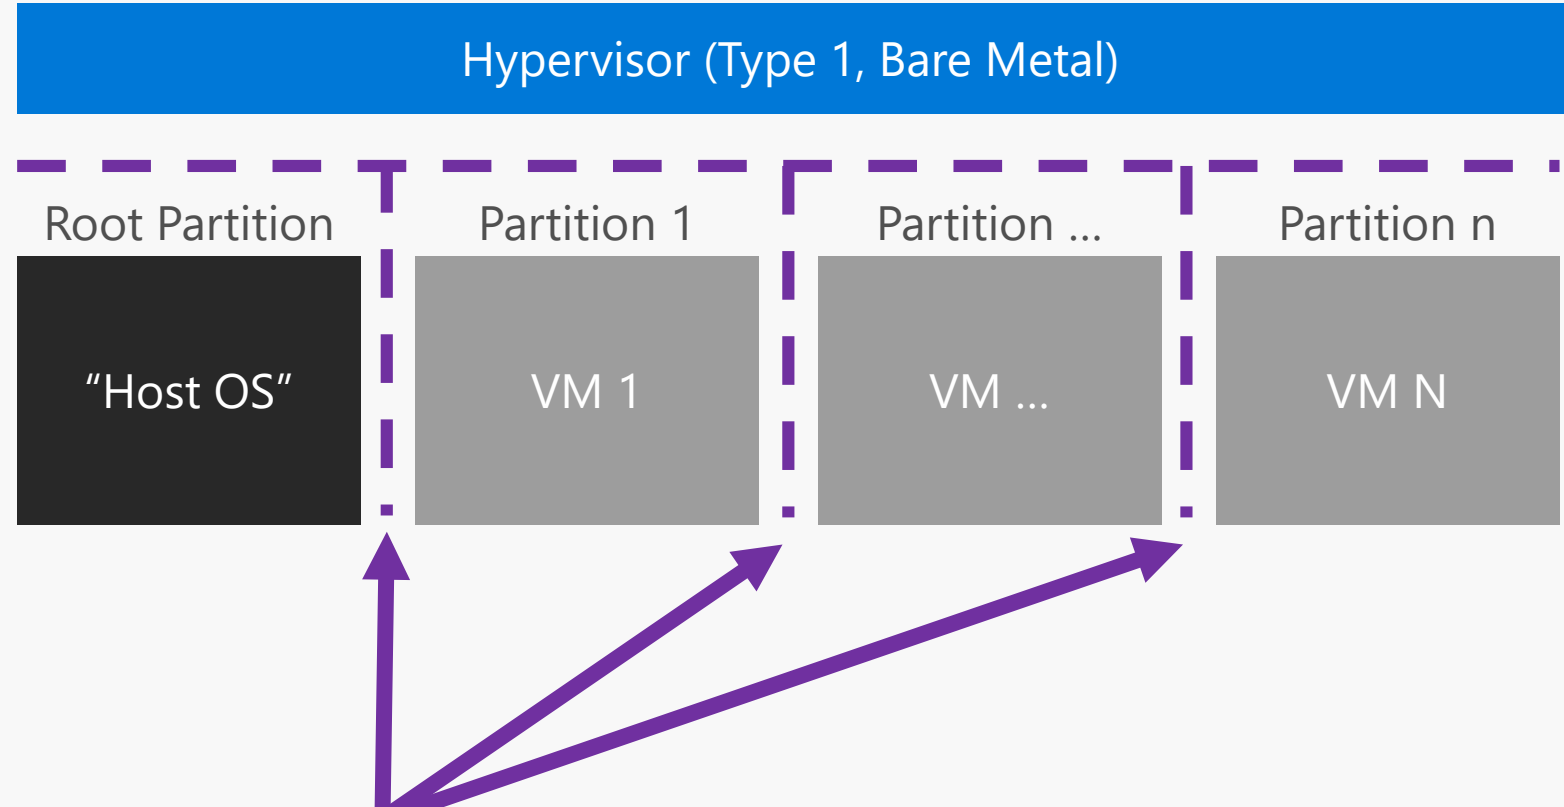
Hyper-V Architecture: Hypervisor

Manages physical address space of partitions (via EPT)

Manages virtualization specific hardware configuration

Handles intercepts (i.e. HyperCall, in/out instructions, CPUID instruction, EPT page fault, etc.)

Interrupt delivery to guests



Hypervisor EPT enforces physical memory isolation between partitions

Most Hyper-V attack surface is not in the hypervisor

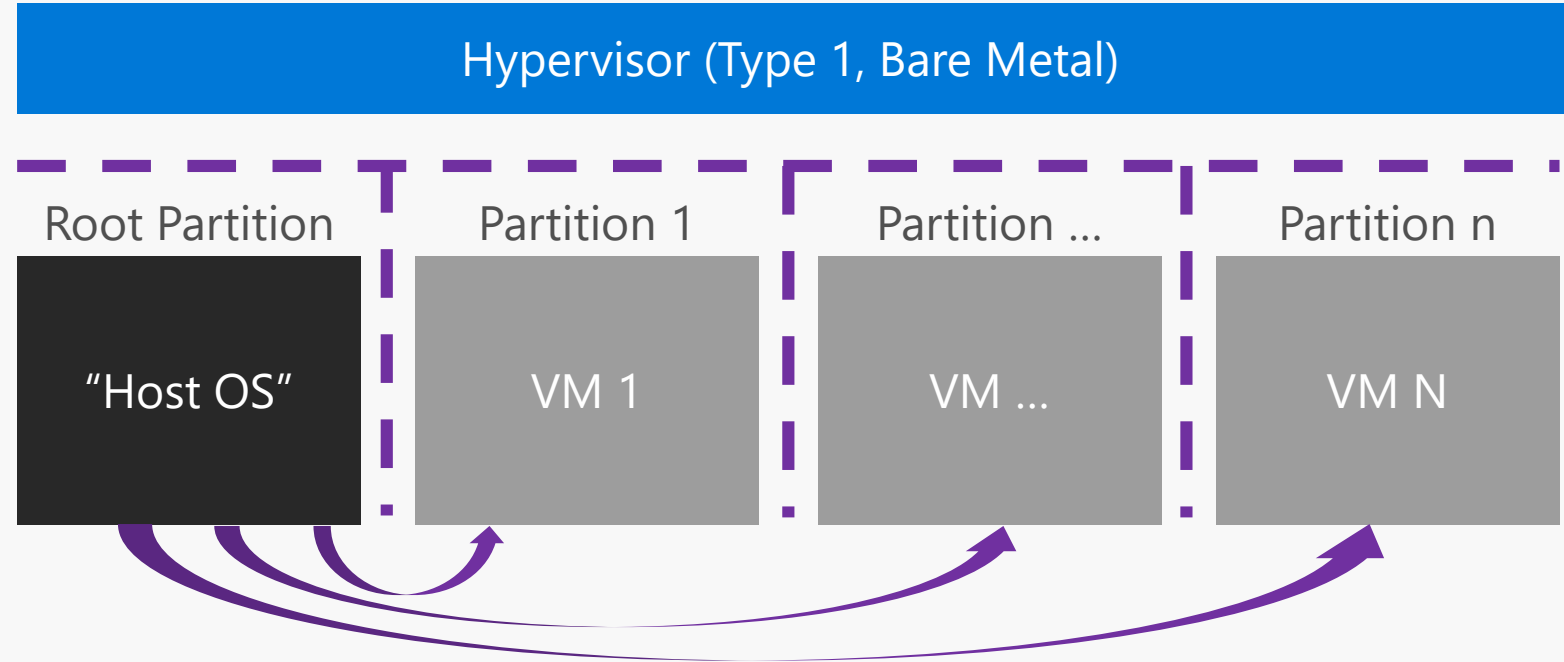
Hyper-V Architecture: Root Partition

Manages other VM's
(create/destroy/etc.)

Access to the physical memory of
other partitions

Access to all hardware

Provides services such as device
emulation, para-virtualized
networking/storage, etc.



Root partition can access other partitions' physical memory

Most Hyper-V attack surface is in the root partition

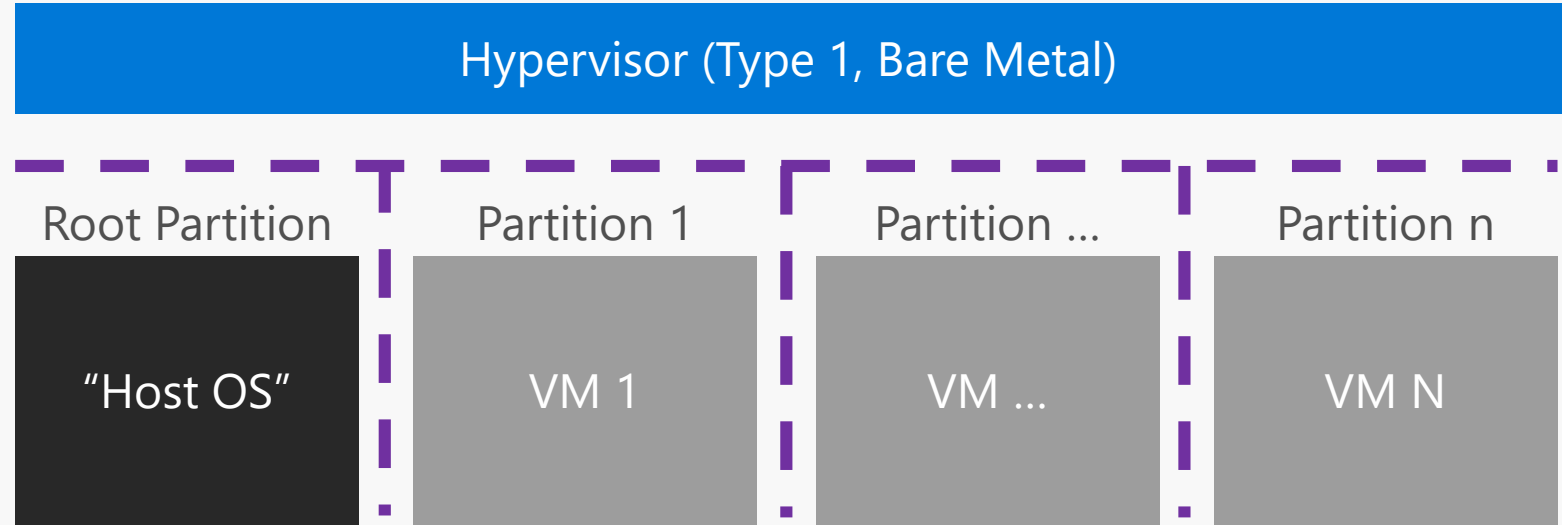
Hyper-V Architecture: Guest Partitions

No access to other partitions
physical memory

No access to hardware

Access to limited set of HyperCalls
(example: faster TLB flush)

No ability to communicate with
partitions other than the root



Communicates with root partition & hypervisor using well defined interfaces

There is no direct guest-to-guest attack surface

Terminology – Physical Memory

- System Physical Address (SPA) – The real physical address.
- Guest Physical Address (GPA) – The physical address a guest sees.
- Guest Physical Address Descriptor List (GPADL) – Conceptually an MDL of GPA's.

Terminology – Types of Components

- Virtual Device (VDEV) – Either an emulated or paravirtualized device hosted in user-mode.
- Virtualization Service Provider (VSP) – Paravirtualized device hosted in kernel. Has an associated VDEV.
- Integration Component (IC) – The same as a VDEV from an attackers POV, user-mode component that guest can communicate with.

Hyper-V Architecture: Root Partition Services

Emulated

Networking (VDEV)
Storage (VDEV)
Floppy Drive (VDEV)
Video (VDEV)
PCI/ISA Bus (VDEV)
Motherboard (VDEV)
Serial Port (VDEV)

Etc...

Para-virtualized

Networking (VSP)
Storage (VSP)
Video (VDEV)
PCI (VSP)

Other

BIOS Firmware
Live Migration
Dynamic Memory
Time sync (IC)
Heartbeat (IC)
SMB Server (VDEV)
Plan9FS (VDEV)

Too much to list...

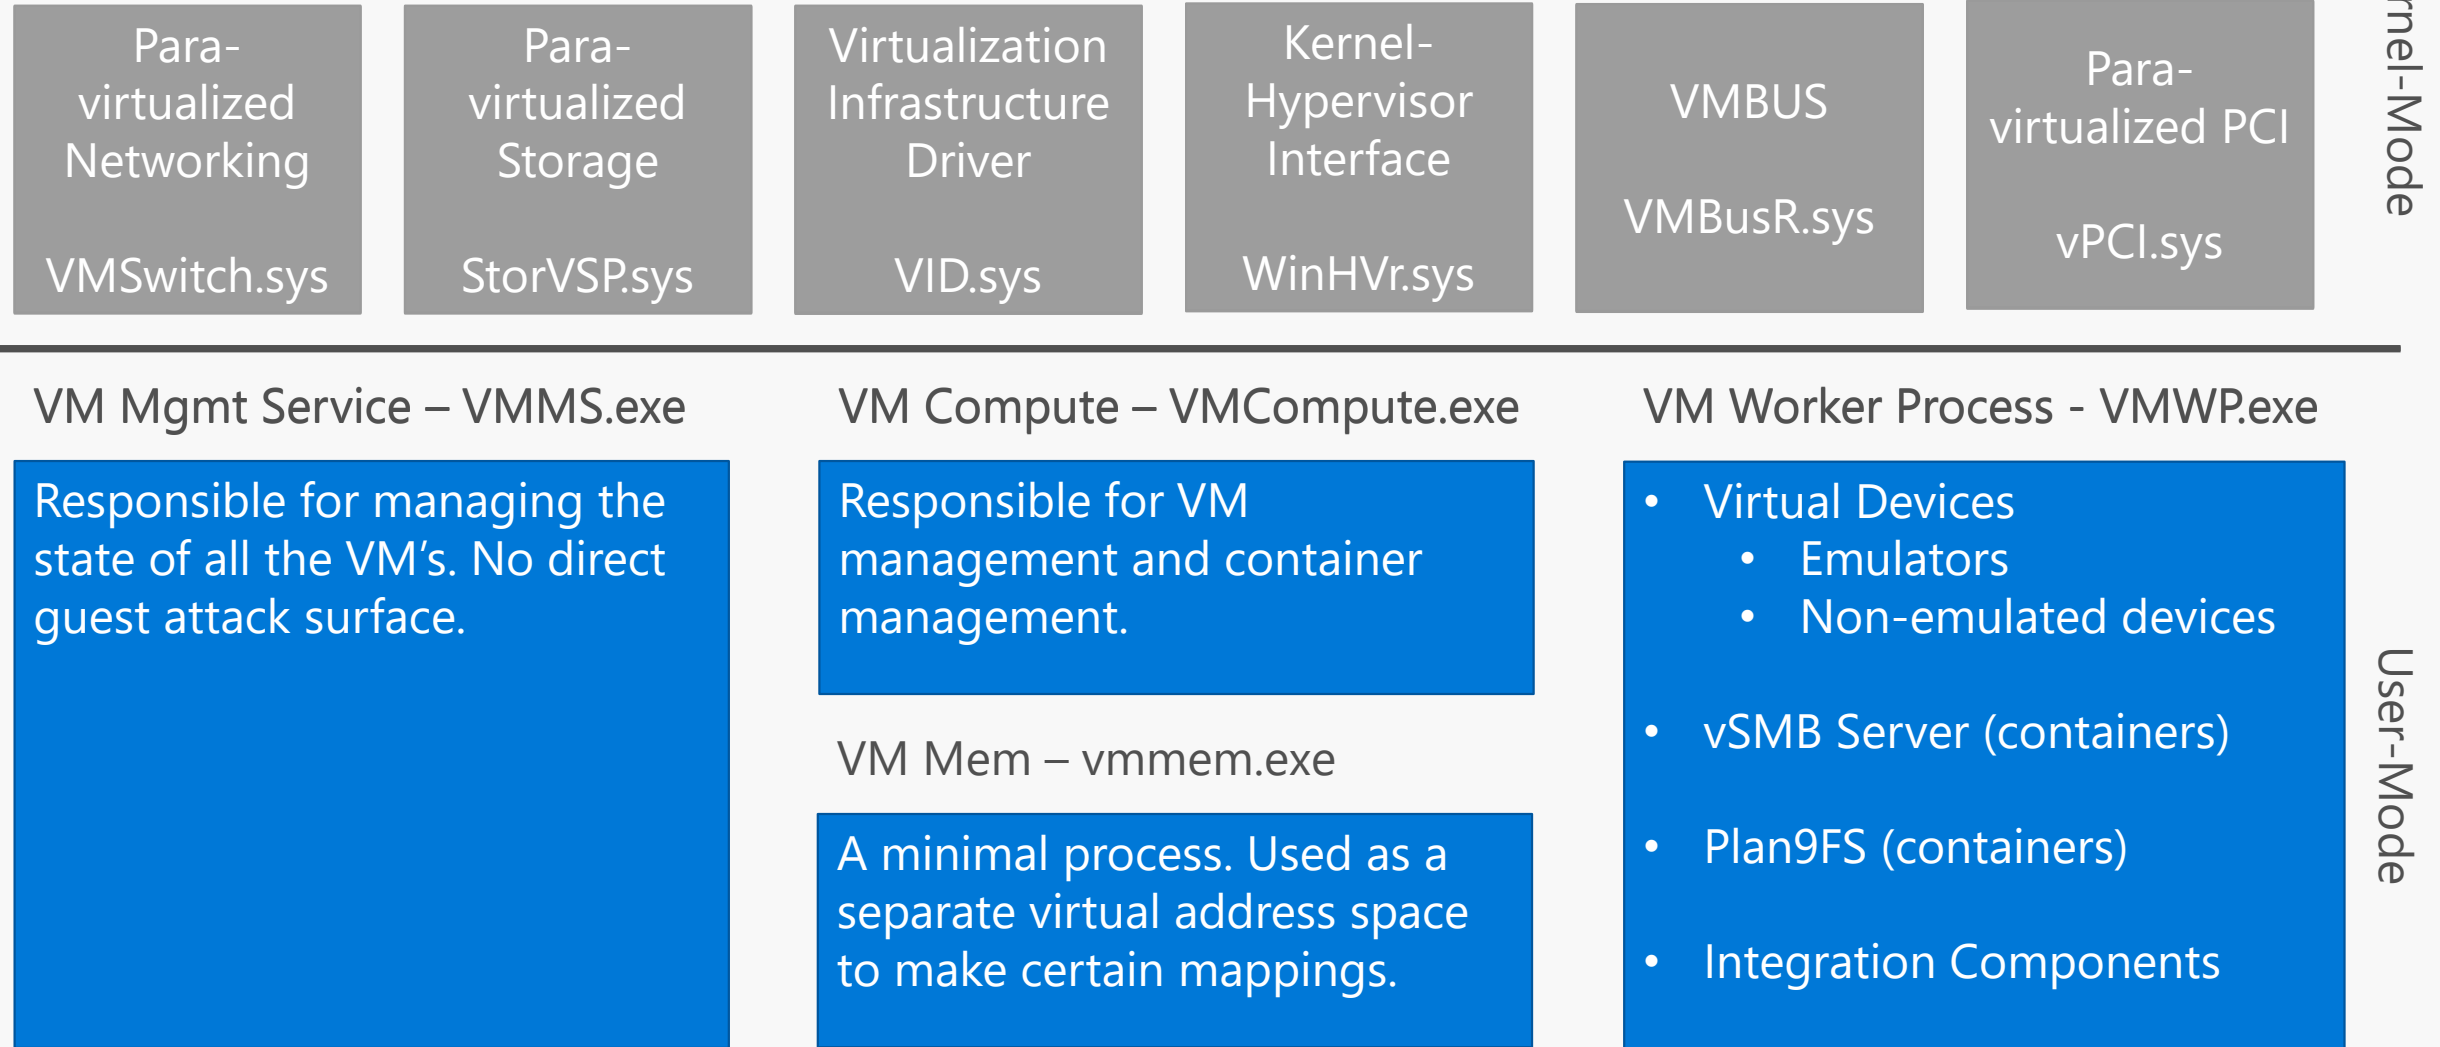
Generation 2 VMs require fewer emulated devices (compared to Generation 1)

Some services mandatory, others configurable

Hyper-V is designed with the principle of least privilege.

As little code as possible is in the hypervisor and root partition kernel.

Hyper-V Architecture: Root Partition



Source code for the guest-side of these VDEV/IC/VSP is in the Linux source tree

Communication Channels (Hypervisor)

Hypercalls

- “System calls” of the hypervisor
- Guest accessible hypercalls are documented as part of the Hyper-V TLFS
- Some Hypercalls pass arguments via registers, others use physical pages (GPA in register)

Faults

- Triple fault, EPT page faults (i.e. permission faults, GPA not mapped, etc.)
- This is how MMIO can be virtualized by VDEV's (fault on access to virtual MMIO range)

Instruction Emulation

- Attempt to execute instructions such as CPUID, RDTSC, RDPMC, INVLPG, IN, OUT, etc.

Register Access

- Attempt to read/write control registers, MSR's

Overlay Pages

- A way for the hypervisor to forcibly map a physical page in to a partition
- Example: Hypercall code page
- Primarily used to communicate data to a guest partition

Communication Channels (Kernel-Mode)

VMBUS

- High-speed communication channel accessed through via Kernel Mode Client Library (KMCL) abstraction layer

Extended Hypercalls

- Hypercalls that the hypervisor forwards directly to the VID
- Very few

Aperture

- Host can map guest physical memory and interact with it
- Rarely used by kernel

Intercept Handling

- Hypervisor forwards some intercepts it receives to the host for processing
 - IO port read/write (does it need emulation?)
 - EPT faults: is the memory paged out?, is that memory a virtual MMIO page?
 - Etc.

Communication Channels (User-Mode)

IO Ports

- User-mode components can register for notifications when particular IO ports are written/read
- Used to emulate hardware

MMIO

- Components can register GPA ranges as MMIO ranges, receive notifications when the ranges are written/read
- Used to emulate hardware

VMBUS

- High-speed communication channel accessed through named pipes or sockets

Aperture


- Map guest physical addresses into the virtual address space of VMWP
- Need to be careful to avoid shared-memory issues such as double-fetch

Read/Write Notifications

- Triggered when a specified GPA is read/written, EIP is not advanced (no emulation)
- Used to track when pages are dirtied while live migrating (as an example)

VMBUS

Shared memory (ring buffer) based communication channel between guest and host



Components interact with VMBUS through abstraction layers

Linux Integration Drivers implement the protocol, good for reverse engineering

VMBUS – KMCL

- Used by VSP's (VMSwitch, StorVSP, vPCI)
- Built around callbacks (i.e. callback on message receive)
 - Callbacks for other events such as channel closure, message sent complete, etc.
- Message received gets copied to non-shared memory
- “External Data” – A GPADL attached to a message which describes guest physical addresses containing additional message data
 - Must be mapped explicitly as an MDL
 - Must be accessed carefully, physical pages are also mapped in guest read/write

KMCL - Packet Receive Entry Point

```
VmbChannelInitSetProcessPacketCallbacks(
```

```
    _In_ VMBCHANNEL Channel,
```

```
    _In_ PFN_VMB_CHANNEL_PROCESS_PACKET ProcessPacketCallback,
```

```
    _In_opt_ PFN_VMB_CHANNEL_PROCESSING_COMPLETE ProcessingCompleteCallback
```

```
)
```

Called to process each packet received from the guest

```
VOID
```

```
EVT_VMB_CHANNEL_PROCESS_PACKET(
```

```
    _In_ VMBCHANNEL Channel,
```

```
    _In_ VMBPACKETCOMPLETION Packet,
```

```
    _In_reads_bytes_(BufferLength) PVOID Buffer,
```

```
    _In_ UINT32 BufferLength,
```

```
    _In_ UINT32 Flags
```

```
);
```

Calls to this function are serialized per-channel

Called after a group of packets has been delivered

```
VOID
```

```
EVT_VMB_CHANNEL_PROCESSING_COMPLETE(
```

```
    _In_ VMBCHANNEL Channel,
```

```
    _In_ UINT32 PacketsProcessed
```

```
);
```

Buffer contains guest-controlled data, NOT in shared memory

VMBUS - Pipes

- Most common VMBUS interface used by user-mode
- Component makes channel offer to guest, receives handle to VMBUS pipe
 - VmBusPipeServerOfferChannel
 - VmBusPipeServerOfferChannelEx
 - Or via wrapper such as VmBusPipeIO class (which uses the above mechanisms)
- Interaction
 - ReadFile/WriteFile
 - IO Completion (asynchronous)
 - Commonly registered with VmCompletionHandlerIo::AssociateHandle (CreateThreadpoolIo)
 - IO completions commonly delivered to: VmNewThreadpool::IoCompletionCallback

IO Port / MMIO Entry Points

IO port being read/written

Size can be: 1, 2, 4

Data (stored in UINT32)

```
HRESULT NotifyMmioRead(  
    [in]          UINT64 RangeBase,  
    [in]          UINT64 RangeOffset,  
    [in]          UINT64 NumberOfBytes,  
    [out, size_is(NumberOfBytes)] BYTE ReadBuffer[] );
```

```
HRESULT NotifyMmioWrite(  
    [in]          UINT64 RangeBase,  
    [in]          UINT64 RangeOffset,  
    [in]          UINT64 NumberOfBytes,  
    [in, size_is(NumberOfBytes)] const BYTE WriteBuffer[] );
```

```
HRESULT NotifyIoPortRead(  
    [in]  VID_IO_PORT_ADDRESS IoAddress,  
    [in]  UINT16 AccessSize,  
    [out] UINT32* ReadData );
```

```
HRESULT NotifyIoPortWrite(  
    [in] VID_IO_PORT_ADDRESS IoAddress,  
    [in] UINT16 AccessSize,  
    [in] UINT32 WriteData );
```

Base MMIO range

Offset into MMIO range

Size of MMIO access

Read/write buffer



Finding bugs!

Note: The vulnerabilities discussed in the following slides have been resolved

A word on symbols...

Virtualization Blog

Information and announcements from Program Managers, Product Managers, Developers and Testers in the Microsoft Virtualization team.

Hyper-V symbols for debugging

★★★★★

April 25, 2018 by Lars Iwer [MSFT] // 0 Comments

 Share 9  26  0

Having access to debugging symbols can be very handy, for example when you are

- A partner building solutions leveraging Hyper-V,
- Trying to debug a specific issue, or
- **Searching for bugs to participate in the Microsoft Hyper-V Bounty Program.**

Starting with symbols for Windows Server 2016 with an installed April 2018 cumulative update, we are now providing access to most Hyper-V-related symbols through the public symbol servers. Here are some of the symbols that are available right now:

```
SYMCHK: vmbuspipe.dll [10.0.14393.2007 ] PASSED - PDB: vmbuspipe.pdb DBG:
SYMCHK: vmbuspiper.dll [10.0.14393.2007 ] PASSED - PDB: vmbuspiper.pdb DBG:
SYMCHK: vmbusvdev.dll [10.0.14393.2007 ] PASSED - PDB: vmbusvdev.pdb DBG:
SYMCHK: vmchipset.dll [10.0.14393.2007 ] PASSED - PDB: VmChipset.pdb DBG:
SYMCHK: vmcompute.dll [10.0.14393.2214 ] PASSED - PDB: vmcompute.pdb DBG:
```

- More details at <https://blogs.technet.microsoft.com/virtualization/2018/04/25/hyper-v-symbols-for-debugging/>

Vulnerabilities

- VMBUS induced vulnerabilities

CVE-2017-0051 - VMSwitch VmsMpCommonPvtSetNetworkAddress Out-of-Bounds Read Vulnerability

CVE-2018-0964 - vPCI VpciMsgCreateInterruptMessage Uninitialized Stack Object

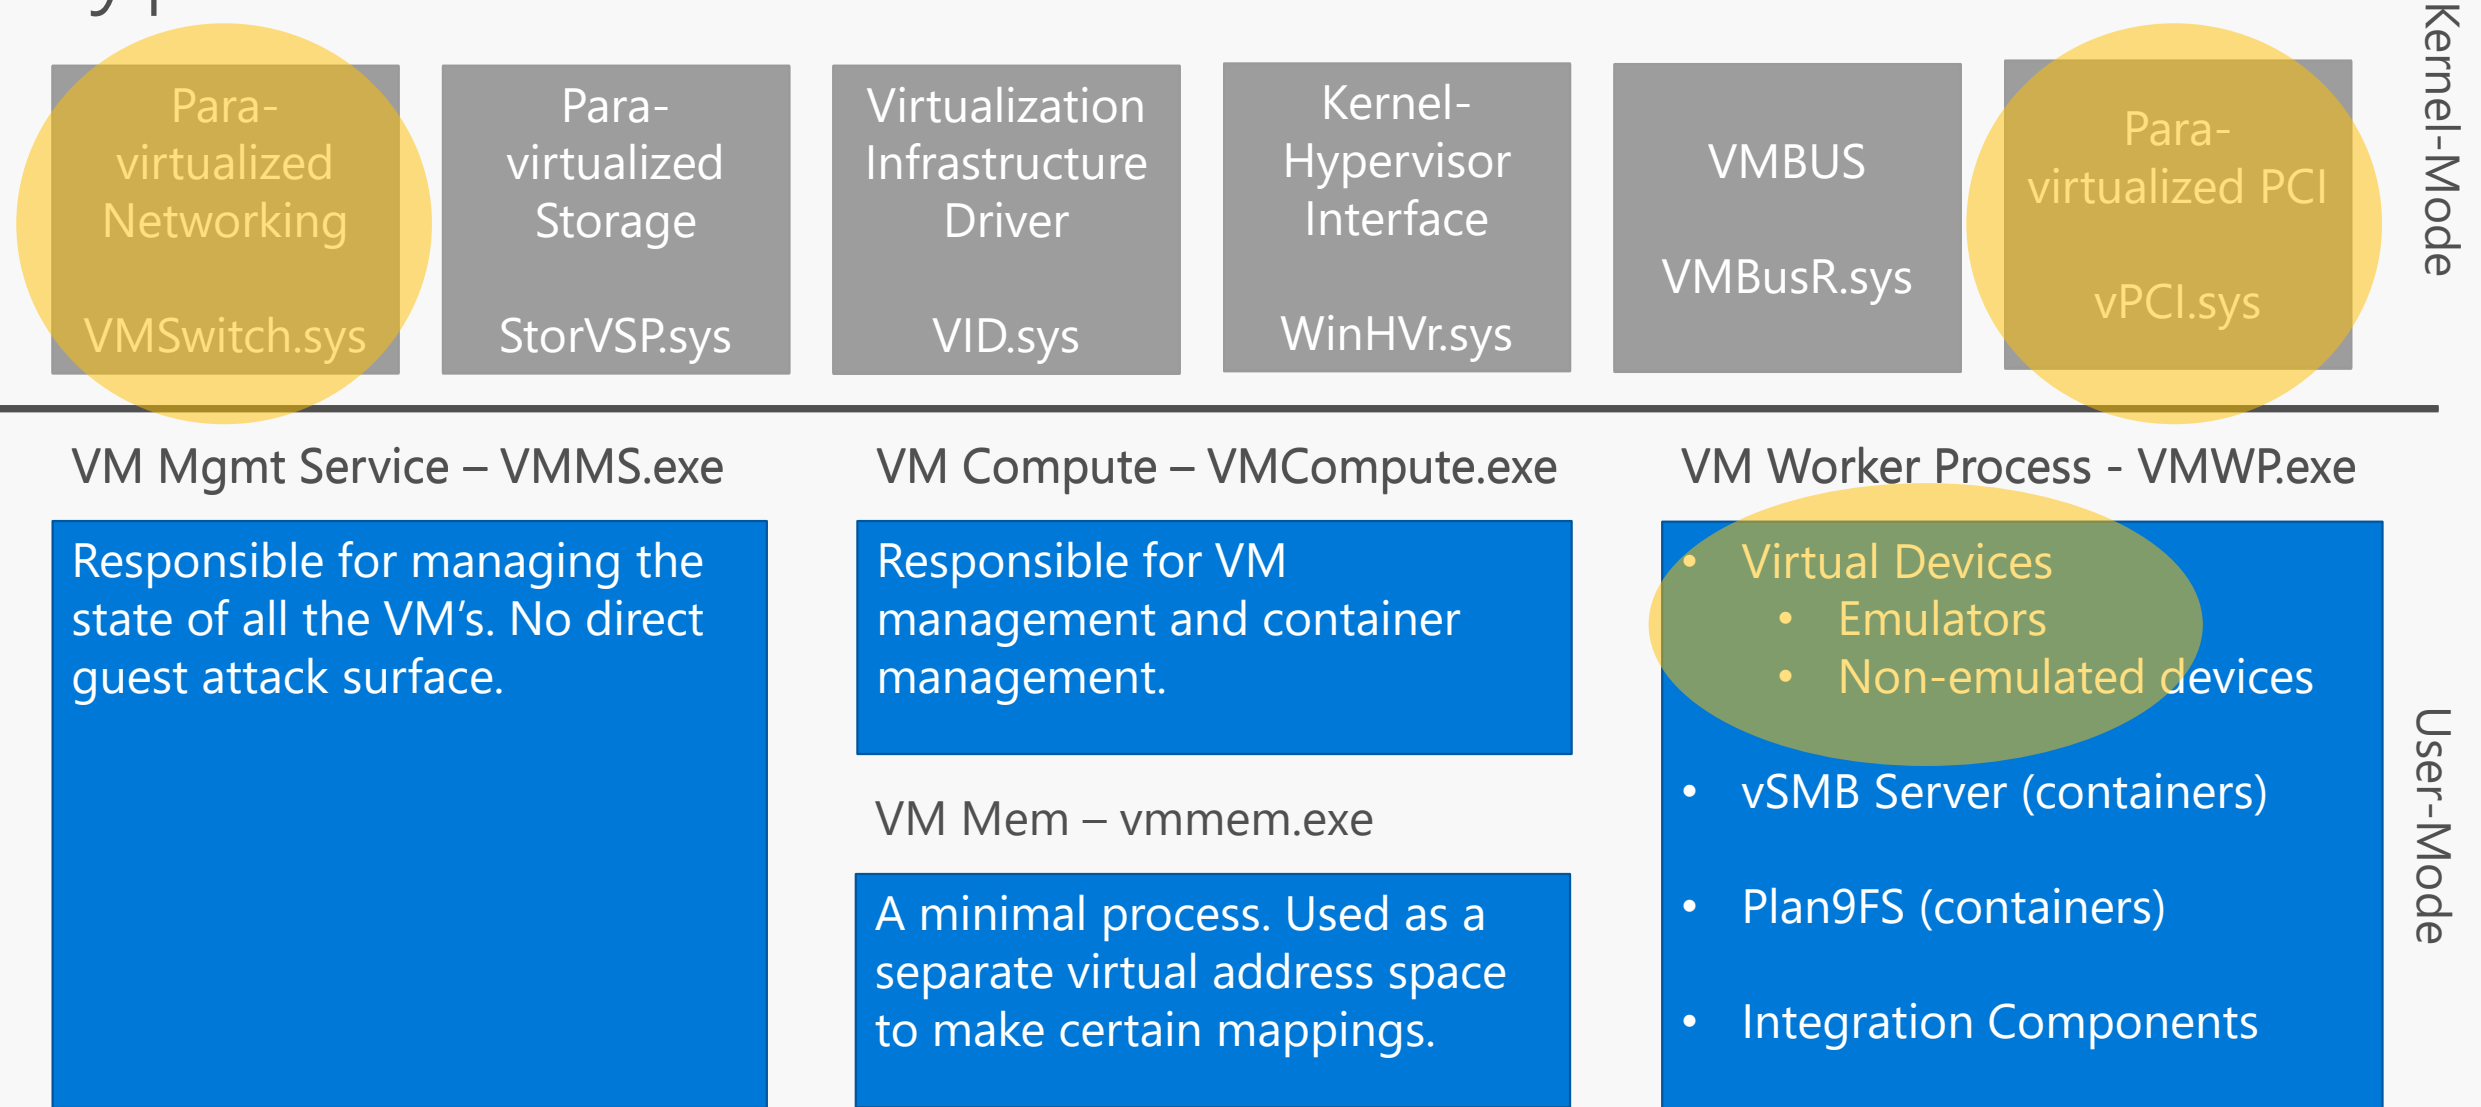
CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field

- Intercepted I/O vulnerabilities

CVE-2018-0888 - Information disclosure during MMIO emulation

CVE-2018-0959 - Out-of-Bounds Read/Write in VmEmulatedStorage

Hyper-V Architecture: Root Partition



Source code for the guest-side of these VDEV/IC/VSP is in the Linux source tree

Vulnerabilities

- VMBUS induced vulnerabilities

`CVE-2017-0051 - VMSwitch VmMpCommonPvtSetNetworkAddress Out-of-Bounds Read Vulnerability`

`CVE-2018-0964 - vPCI VpciMsgCreateInterruptMessage Uninitialized Stack Object`

`CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field`

- Intercepted I/O vulnerabilities

`CVE-2018-0888 - Information disclosure during MMIO emulation`

`CVE-2018-0959 - Out-of-Bounds Read/Write in VmEmulatedStorage`

CVE-2017-0051 – VMSwitch VmsMpCommonPvtSetNetworkAddress Out-of-Bounds Read Vulnerability

- Found by Peter Hlavaty (Tencent)
- Issue introduced in RS1
- In error paths, VmsMpCommonPvtSetNetworkAddress passes an attacker controlled WSTR to a logging function
 - Attacker may not null-terminate this WSTR
 - Error logging function looks for null, can read out-of-bounds until page fault
- **Host DoS from the guest**
- **Hyper-V Bug Bounty today: \$15,000**

CVE-2017-0051 - VMSwitch VmsMpCommonPvtSetNetworkAddress Out-of-Bounds Read Vulnerability

```
70 int qilin2hyperv_ddos(  
71     struct rndis_device *rdev  
72 )  
73 {  
74     struct rndis_config_parameter_info *cpi;  
75     wchar_t *cfg_nwadr, *cfg_mac;  
76     struct rndis_set_request *set;  
77     struct rndis_request* request;  
78  
79     u32 extlen = sizeof(struct rndis_config_parameter_info) + 0x40;  
80  
81     request = get_rndis_request(rdev, RNDIS_MSG_SET,  
82         RNDIS_MESSAGE_SIZE(struct rndis_set_request) + extlen);  
83     if (!request)  
84         return -ENOMEM;  
85  
86     memset(cpi, 'A', set->info_buflen);  
87  
88     cpi->parameter_name_offset =  
89         sizeof(struct rndis_config_parameter_info);  
90     /* Multiply by 2 because host needs 2 bytes (utf16) for each  
91     cpi->parameter_name_length = 2*NWADR_STRLEN;  
92     cpi->parameter_type = RNDIS_CONFIG_PARAM_TYPE_STRING;  
93     cpi->parameter_value_offset = extlen - 2;  
94     /* Multiply by 4 because each MAC byte displayed as 2 utf16 chars */  
95     cpi->parameter_value_length = 2;  
96  
97     cfg_nwadr = (wchar_t *) ((ulong)cpi + cpi->parameter_name_offset);  
98     cfg_mac = (wchar_t *) ((ulong)cpi + cpi->parameter_value_offset);  
99     utf8s_to_utf16s(NWADR_STR, NWADR_STRLEN, UTF16_HOST_ENDIAN,  
100         cfg_nwadr, NWADR_STRLEN);  
101  
102     return rndis_filter_send_request(rdev, request);  
103 }  
104 }
```

Patch the Linux
drivers in
rndis_filter.c

Cause an error to log
the long string

Run ifconfig

VmsMpCommonPvtSetNetwork
Address with a long
unterminated string

RNDIS packet sent
to the VMBUS

- How is the RNDIS packet processed?

```
00 vmswitch!RndisDevHostQueueWorkItem  
01 vmswitch!RndisDevHostDispatchControlMessage  
02 vmswitch!VmsVmNicPvtKmclProcessingComplete  
03 vmswitch!VmsVmNicPvtKmclProcessPacket
```

RndisDevHostQueueWorkItem proc near

```
sub     rsp, 28h  
xor     eax, eax  
lea     r8d, [rax+1]  
lock cpxchg [rcx+98h], r8d  
jnz     short loc_1C001E4AC  
lock add [rcx+0A0h], r8d  
mov     r9, rcx  
lea     rdx, RndisDevHostControlMessageWorkerRoutine  
mov     rcx, [rcx+90h]  
call    cs:__imp_IoQueueWorkItemEx
```

```
0:003> kc 10  
# Call Site  
00 nt!???:FNODOBFM::string'  
01 nt!MmAccessFault  
02 nt!KiPageFault  
03 vmswitch!WPP_RECORDER_SF_qSd  
04 vmswitch!VmsMpCommonPvtSetNetworkAddress  
05 vmswitch!VmsMpCommonPvtSetRequestCommon  
06 vmswitch!VmsMpCommonSetRequest  
07 vmswitch!VmsVmNicPvtRdisDeviceSetRequest  
08 vmswitch!RdisDevHostHandleSetMessage  
09 vmswitch!RdisDevHostControlMessageWorkerRoutine  
0a nt!IopProcessWorkItem  
0b nt!ExpWorkerThread  
0c nt!PspSystemThreadStartup  
0d nt!KiStartSystemThread
```

From receiving the packet to VmsMpCommonPvtSetNetworkAddress

Other VMSwitch issues

- Kostya Kortchinsky (Google):
 - <https://bugs.chromium.org/p/project-zero/issues/detail?id=688>
 - <https://bugs.chromium.org/p/project-zero/issues/detail?id=689>
 - <https://bugs.chromium.org/p/project-zero/issues/detail?id=690>
- MS17-008
 - **Jordan Rabet's talk at Black Hat**
<https://www.blackhat.com/us-18/briefings.html#hardening-hyper-v-through-offensive-security-research>

Vulnerabilities

- VMBUS induced vulnerabilities

CVE-2017-0051 - VMSwitch VmsMpCommonPvtSetNetworkAddress Out-of-Bounds Read Vulnerability

CVE-2018-0964 - vPCI VpciMsgCreateInterruptMessage Uninitialized Stack Object

CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field


- Intercepted I/O vulnerabilities

CVE-2018-0888 - Information disclosure during MMIO emulation

CVE-2018-0959 - Out-of-Bounds Read/Write in VmEmulatedStorage

CVE-2018-0964 – vPCI VpciMsgCreateInterruptMessage Uninitialized Stack Object

- Found by the Virtualization Security Team (Microsoft)
- VirtualBusChannelProcessPacket in vpcivsp.sys, switch of 25 cases:



- VirtualDeviceCreateSingleInterrupt doesn't always initialize TranslatedMessage

```
typedef struct _VPCI_MESSAGE_RESOURCE_2
{
    union
    {
        struct
        {
            USHORT    Reserved;
            USHORT    MessageCount;
            ULONG      DataPayload;
            ULONG64    Address;
            USHORT     Reserved2[27];
        } Remapped;
    }
};
```

```
status = VirtualDeviceCreateSingleInterrupt(device,
                                             &transCreateIntPacket2,
                                             &TranslatedMessage
                                             );

RtlSecureZeroMemory(&createIntReply, sizeof(createIntReply));

createIntReply.ReplyHeader.Status = status;
createIntReply.TranslatedMessage.Remapped.Reserved = TranslatedMessage.Remapped.Reserved;
createIntReply.TranslatedMessage.Remapped.MessageCount = TranslatedMessage.Remapped.MessageCount;
createIntReply.TranslatedMessage.Remapped.DataPayload = TranslatedMessage.Remapped.DataPayload;
createIntReply.TranslatedMessage.Remapped.Address = TranslatedMessage.Remapped.Address;

VirtualBusPacketComplete(device->VirtualBus,
                        PacketCompletionContext,
                        &createIntReply,
                        sizeof(createIntReply));
```

CVE-2018-0964 - vPCI VpciMsgCreateInterruptMessage Uninitialized Stack Object

- How to reach that code?
- Look for xrefs to VmbChannelSendSynchronousRequest or VmbPacketSend in vpci.sys in the guest
- Break on FdoProtocolCommunication to see the handshake on the VMBUS
- Replay your own packets

VpciMsgCreateInterruptMessage = 0x42490014

- **Leak sensitive information from the host kernel**
- **Hyper-V Bug Bounty today: \$25,000**

VpciMsgQueryProtocolVersion
= 0x42490013



```
00000001C000BB8A loc_1C00BB8A:          CODE XREF: FdoProtocolCommunication+E74j
00000001C000BB8A          mov     eax, [r14]
00000001C000BB8D          mov     [rsp+68h+arg_14], eax
00000001C000BB94          mov     [rsp+68h+arg_10], 42490013h
00000001C000BB9F          mov     rcx, cs:WPP_GLOBAL_Control ; __annotation("TMF:",
                                ; "457ffa6b-7a75-3e8b-0f99-c3feedc37640 :
                                ; "#typev Unknown_cxx00 18 \"%0%10!p!: Ser
                                ; "{", "Arg, ItemPtr -- 10", "Arg, ItemL
                                ; "PUBLIC_TMF:")
00000001C000BBA6          mov     r9d, 12h ; id
00000001C000BBAC          mov     [rsp+68h+_a2], eax ; _a2
00000001C000BBB0          mov     dl, 4 ; level
00000001C000BBB2          mov     [rsp+68h+_a1], rdi ; _a1
00000001C000BBB7          mov     [rsp+68h+traceGuid], rbp ; traceGuid
00000001C000BBBC          mov     rcx, [rcx+40h] ; AutoLogContext
00000001C000BBCE          lea     r8d, [r9-0Ch] ; flags
00000001C000BBCE          call    WPP_RECORDER_SF_qd
00000001C000BBCE          and     [rsp+68h+var_30], 0
00000001C000BBCE          lea     rax, [rsp+68h+arg_8]
00000001C000BBCE          mov     rcx, [rdi+18h]
00000001C000BBCE          lea     rdx, [rsp+68h+arg_10]
00000001C000BBCE          mov     qword ptr [rsp+68h+_a2], rax
00000001C000BBCE          xor     r9d, r9d
00000001C000BBCE          lea     rax, [rsp+68h+arg_18]
00000001C000BBCE          mov     [rsp+68h+arg_8], 8
00000001C000BBCE          mov     [rsp+68h+_a1], rax
00000001C000BBCE          mov     dword ptr [rsp+68h+traceGuid], 1
00000001C000BBCE          lea     r8d, [r9+8]
00000001C000BBCE          call    cs:__imp_VmbChannelSendSynchronousRequest
```

Vulnerabilities

- VMBUS induced vulnerabilities

CVE-2017-0051 - VMSwitch VmsMpCommonPvtSetNetworkAddress Out-of-Bounds Read Vulnerability

CVE-2018-0964 - vPCI VpciMsgCreateInterruptMessage Uninitialized Stack Object

CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field

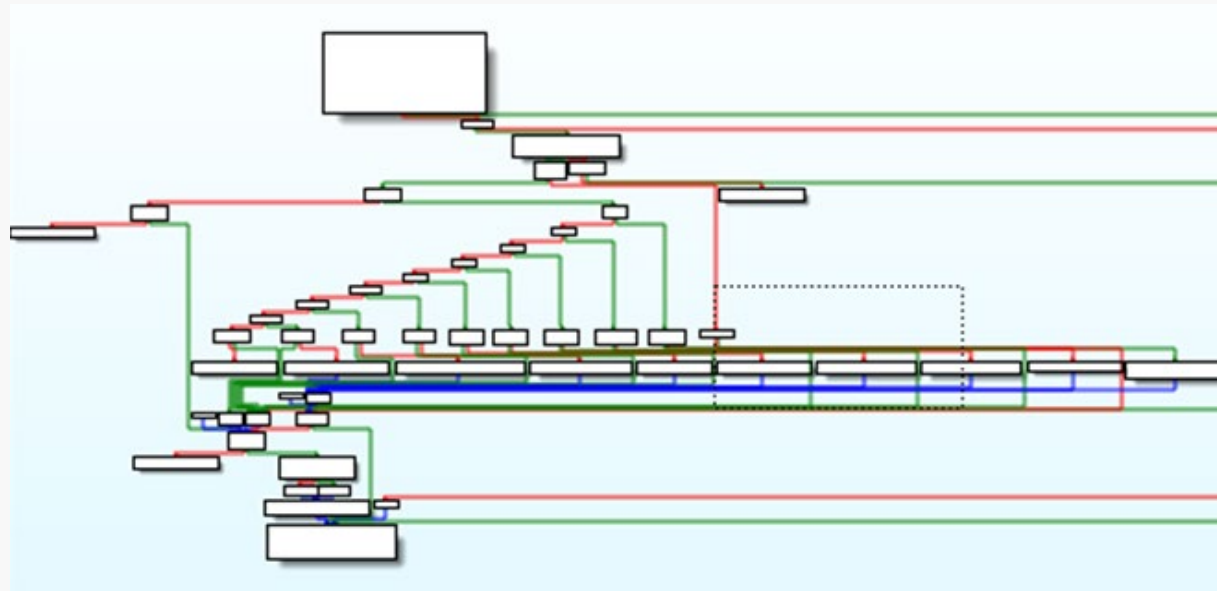
- Intercepted I/O vulnerabilities

CVE-2018-0888 - Information disclosure during MMIO emulation

CVE-2018-0959 - Out-of-Bounds Read/Write in VmEmulatedStorage

CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field

- Found by Nicolas Joly (Microsoft)
- Affects vmwp.exe, relevant code in vmuidevices.dll
- Messages are received by VideoSynthDevice::OnMessageReceived
 - Switch of 9 cases



- Responses are sent by VideoSynthDevice::SendNextMessageInternal
 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse

CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field

```
if (!Accepted)
{
    responseSize = sizeof(SYNTHVID_SUPPORTED_RESOLUTIONS_RESPONSE_MESSAGE);

    response = (PSYNTHVID_SUPPORTED_RESOLUTIONS_RESPONSE_MESSAGE) new(std::nothrow) BYTE[responseSize];
    if (response == NULL)
    {
        hr = E_OUTOFMEMORY;
        goto ErrExit;
    }

    response->Header.Type = SynthvidSupportedResolutionsResponse;
    response->Header.Size = responseSize;
    response->ResolutionCount = 0;
}

hr = SendMessage(&response->Header);
if (FAILED(hr))
```

sizeof(SYNTHVID_SUPPORTED_RES) = 0x8F!

- Leak 0x86 bytes of heap memory to the guest

Hyper-V Bug Bounty Today: \$15,000

- Variant for a stack object in VideoSynthDevice::SendNextMessageInternal

Double your gain with another \$15,000

```
mov     ebp, 8Fh
lea     rdx, std::nothrow_t const std::nothrow ; x
mov     ecx, ebp                ; size
call    operator new[](unsigned __int64,std::nothrow_t const &)
mov     rbx, rax
test    rax, rax
jnz     short loc_18002BE1E
```

```
loc_18002BE1E:
mov     dword ptr [rax], 0Eh
mov     [rax+4], ebp
mov     byte ptr [rax+88h], 0
jmp     loc_18002C1F3
```

```
loc_18002C1F3:                ; Message
mov     rdx, rbx
mov     rcx, rsi                ; this
call    VideoSynthDevice::SendMessageW(SYNTHVID_MESSAGE_HEADER *,bool)
mov     edi, eax
```

Only 9 bytes initialized

CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field

- How to trigger?
 - Relevant code in HyperVideo.sys in the guest
 - Initialization messages sent when the guest loads
 - Break on SynthVidpSendMessageSynchronousLocked
- Example, look at the handshake in SynthVidInitialize:

```
versionRequest->Header.Type = SynthvidVersionRequest;  
versionRequest->Header.Size = sizeof(*versionRequest);  
versionRequest->Version.AsDWORD = SYNTHVID_VERSION_CURRENT;  
  
status = SynthVidpSendMessageSynchronousLocked(  
    libContext,  
    sizeof(*versionRequest),  
    &versionResponse,  
    sizeof(versionResponse),  
    &bytesRead);
```

```
mov     edx, 0Ch          ; SendLength  
lea     r8, [rsp+58h+ReceiveBuffer] ; ReceiveBuffer  
mov     dword ptr [rax], 1  
mov     [rax+4], edx  
lea     r9d, [rdx+2]      ; ReceiveBufferLength  
mov     dword ptr [rax+8], 50003h  
lea     rax, [rsp+58h+BytesRead+28]  
mov     [rsp+58h+BytesRead], rax ; BytesRead  
call    SynthVidpSendMessageSynchronousLocked
```

Change the type, size, content and start fuzzing!

Vulnerabilities

- VMBUS induced vulnerabilities

CVE-2017-0051 - VMSwitch VmsMpCommonPvtSetNetworkAddress Out-of-Bounds Read Vulnerability

CVE-2018-0964 - vPCI VpciMsgCreateInterruptMessage Uninitialized Stack Object

CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field

- Intercepted I/O vulnerabilities

CVE-2018-0888 - Information disclosure during MMIO emulation

CVE-2018-0959 - Out-of-Bounds Read/Write in VmEmulatedStorage

CVE-2018-0888 - Information disclosure during MMIO emulation

- NotifyMmioRead returns "NumberOfBytes" bytes from "ReadBuffer" to the VM
 - Return value is ignored, these bytes are ALWAYS returned to the VM
- If virtual device doesn't populate ReadBuffer, uninitialized stack data is returned to the guest
- This was fixed by initializing ReadBuffer prior to calling NotifyMmioRead
- Found by Joe Bialek (Microsoft)

Hyper-V Bug Bounty Today: \$15,000

```
void BatteryEmulator::NotifyMmioRead(
    _In_ UINT64      RangeBase,
    _In_ UINT64      RangeOffset,
    _In_ UINT64      NumberOfBytes,
    _Out_writes_bytes_(NumberOfBytes) BYTE ReadBuffer[] ) noexcept
{
    if (NumberOfBytes != 4)
        return;
    ...
}
```

Must be initialized by this function

NumberOfBytes != 4 results in ReadBuffer never be initialized

Vulnerabilities

- VMBUS induced vulnerabilities

CVE-2017-0051 - VMSwitch VmsMpCommonPvtSetNetworkAddress Out-of-Bounds Read Vulnerability

CVE-2018-0964 - vPCI VpciMsgCreateInterruptMessage Uninitialized Stack Object

CVE-2017-8706 - VideoSynthDevice::SynthVidSendSupportedResolutionsResponse Uninitialized Object Field

- Intercepted I/O vulnerabilities

CVE-2018-0888 - Information disclosure during MMIO emulation

CVE-2018-0959 - Out-of-Bounds Read/Write in VmEmulatedStorage

CVE-2018-0959 – Out-of-Bounds Read/Write in VmEmulatedStorage

- Anonymously reported
- Affects EmulatedIDE in vmwp.exe, relevant code in VmEmulatedStorage.dll
- Out-of-Bounds Read/Write due to an unexpected internal state and lack of bounds checking in:
 - IdeChannel::ReadDataPort
 - IdeChannel::WriteDataPort

```
UINT8* curBuffer;  
if (Drive.Saved.UseCommandBuffer)  
{  
    curBuffer = (UINT8*)Drive.CommandBuffer;  
}  
else  
{  
    curBuffer = Drive.TrackCacheBuffer + Drive.Saved.DriveStateBufferOffset;  
}
```

DriveStateBufferOffset was not properly set

```
UINT32 curByte = Drive.Saved.CurrentByte;  
UINT32 length = AccessCount * AccessSize;  
  
if (curByte + length > Drive.Saved.TotalBytes)  
{  
    VM_LOG_TRACE(  
        (TraceVDevIdeControllerError,  
        L"[IDE ] Write to data port exceeds TotalBytes."));  
  
    VML_ASSERT(curByte + length <= Drive.Saved.TotalBytes);  
    length = Drive.Saved.TotalBytes - curByte;  
}  
  
// Copy the data.  
RtlCopyMemory(curBuffer + curByte, Buffer, length);  
curByte += length;
```

CVE-2018-0959 – Out-of-Bounds Read/Write in VmEmulatedStorage

- The poc just consists of a series of **out port, value**
- Allows arbitrary Read/Write on a 4GB area

```
(1620.678): Access violation - code c0000005 (first/second chance not available)
ucrtbase!MoveSmall+0x76:
00007ff9`9ad88866 418902          mov     dword ptr [r10],eax ds:00000297`5f670200=????????
0:003> kc 10
# Call Site
00 ucrtbase!MoveSmall
01 VmEmulatedStorage!IdeChannel::WriteDataPort
02 VmEmulatedStorage!IdeChannel::WritePort
03 VmEmulatedStorage!IdeChannel::AltWriteIoPort
04 VmEmulatedStorage!IdeControllerDevice::NotifyIoPortWrite
05 vmwp!VmbCallback::NotifyIoPortWrite
06 vmwp!EmulatorVp::DispatchIoPortOperation
07 vmwp!EmulatorVp::TrySimpleIoEmulation
08 vmwp!EmulatorVp::TryIoEmulation
```

- Found by fuzzing I/O in the Ide Controller with page heap enabled on vmwp.exe
- Top bounty awarded for Hyper-V so far!

★★★★ \$150,000 ★★★★★



TENCENT SECURITY CONFERENCE 2018
2018腾讯安全国际技术峰会

Closing Thoughts

Closing Thoughts

- Hyper-V presents an interesting and well designed target
- Please help us find bugs, we are looking forward to paying a \$250,000 bounty!
- Check out Jordan Rabet's talk on Hyper-V exploitation & mitigations
 - **"HARDENING HYPER-V THROUGH OFFENSIVE SECURITY RESEARCH"**



TENCENT SECURITY CONFERENCE 2018
2018腾讯安全国际技术峰会

THANKS